



# Authenticated Encryption

---

Active attacks on  
CPA-secure encryption

# Recap: the story so far

**Confidentiality:** semantic security against a CPA attack

- Encryption secure against eavesdropping only

**Integrity:**

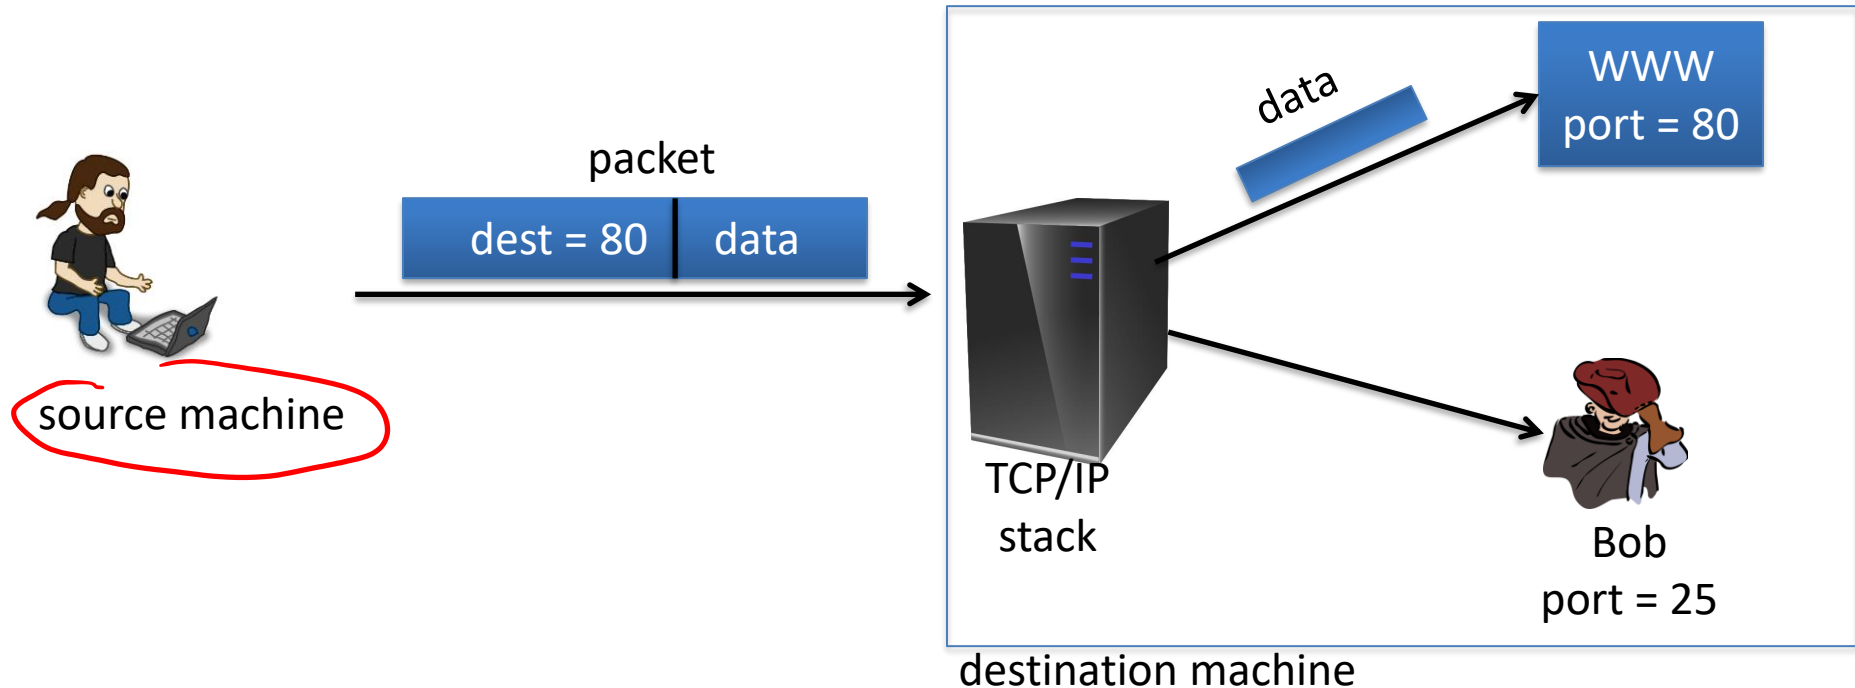
- Existential unforgeability under a chosen message attack
- CBC-MAC, HMAC, PMAC, CW-MAC

This module: encryption secure against tampering (*active adversary*)

- Ensuring both confidentiality and integrity

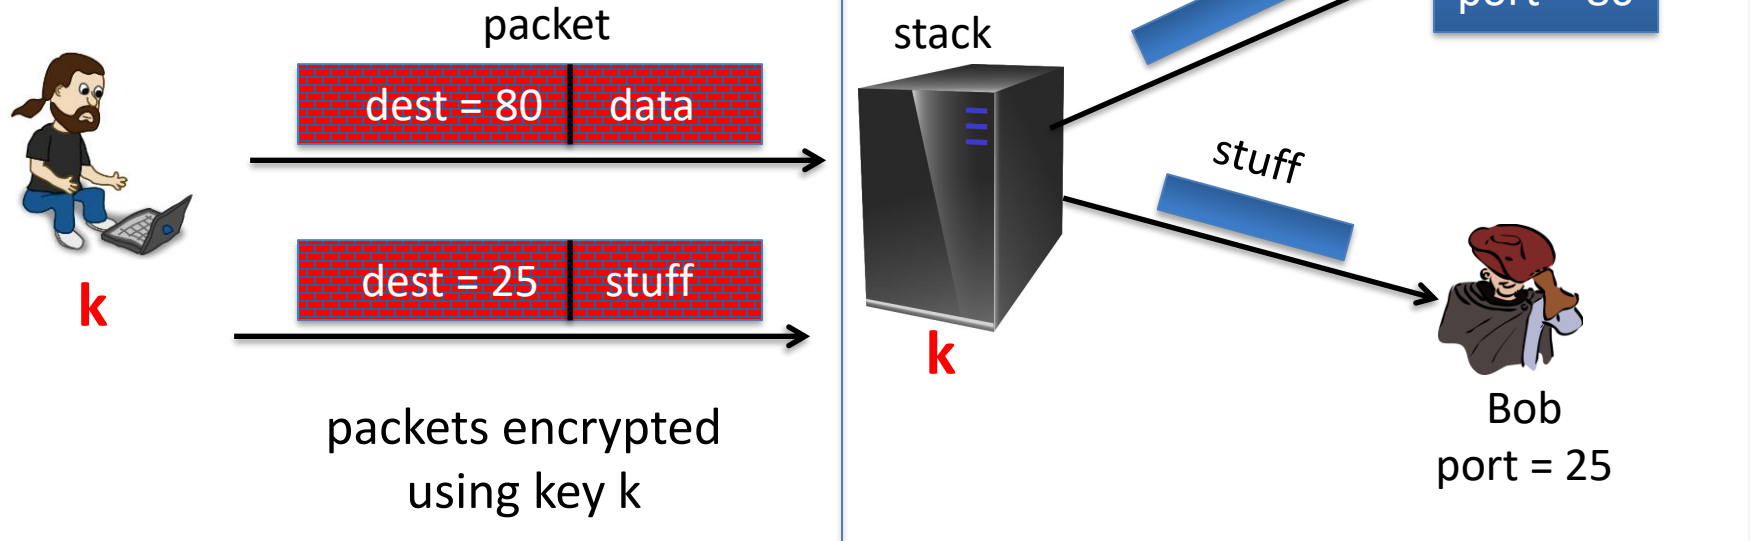
# Sample tampering attacks

TCP/IP: (highly abstracted)




# Sample tampering attacks

IPsec: (highly abstracted)




# Reading someone else's data

Note: attacker obtains decryption of any ciphertext beginning with “dest=25”

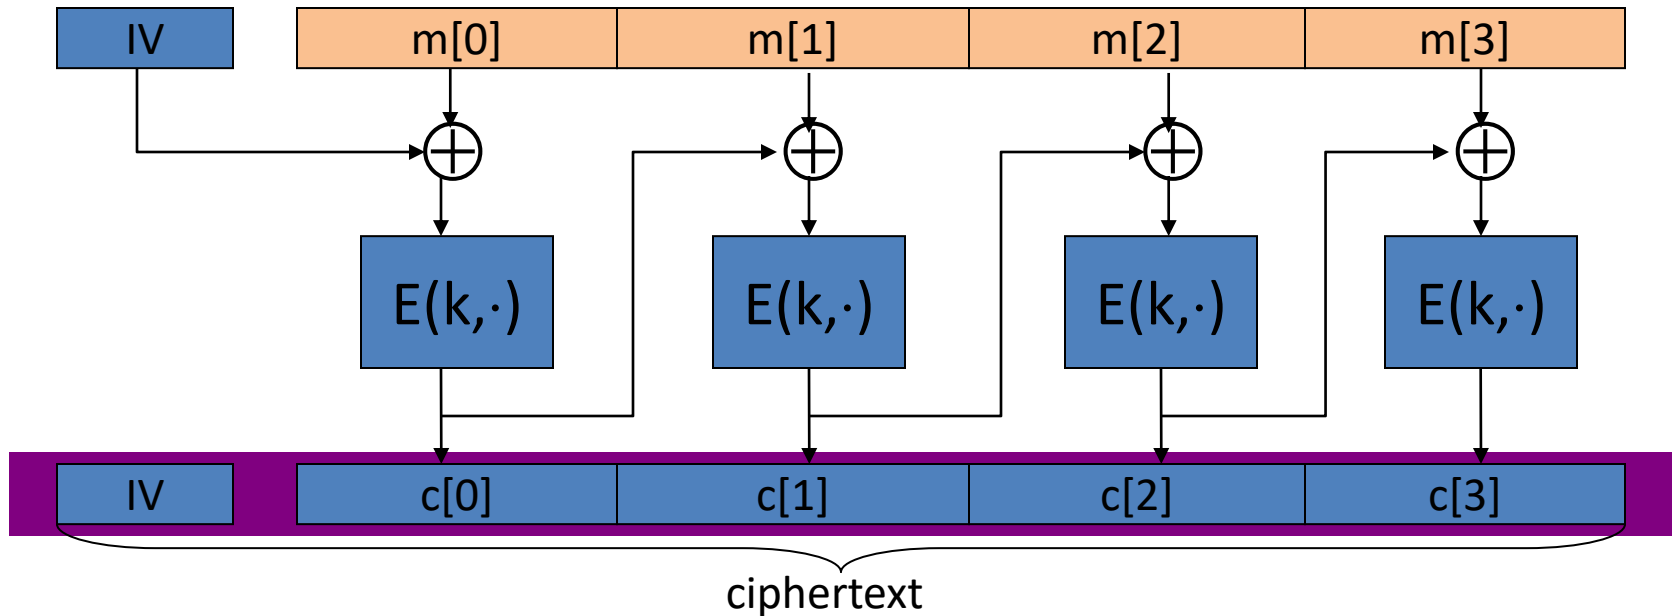


Easy to do for CBC with rand. IV  
(only IV is changed)



# Review: CBC with random IV

Let  $(E,D)$  be a PRP.  $E_{\text{CBC}}(k,m)$ : choose random  $IV \in X$  and do:





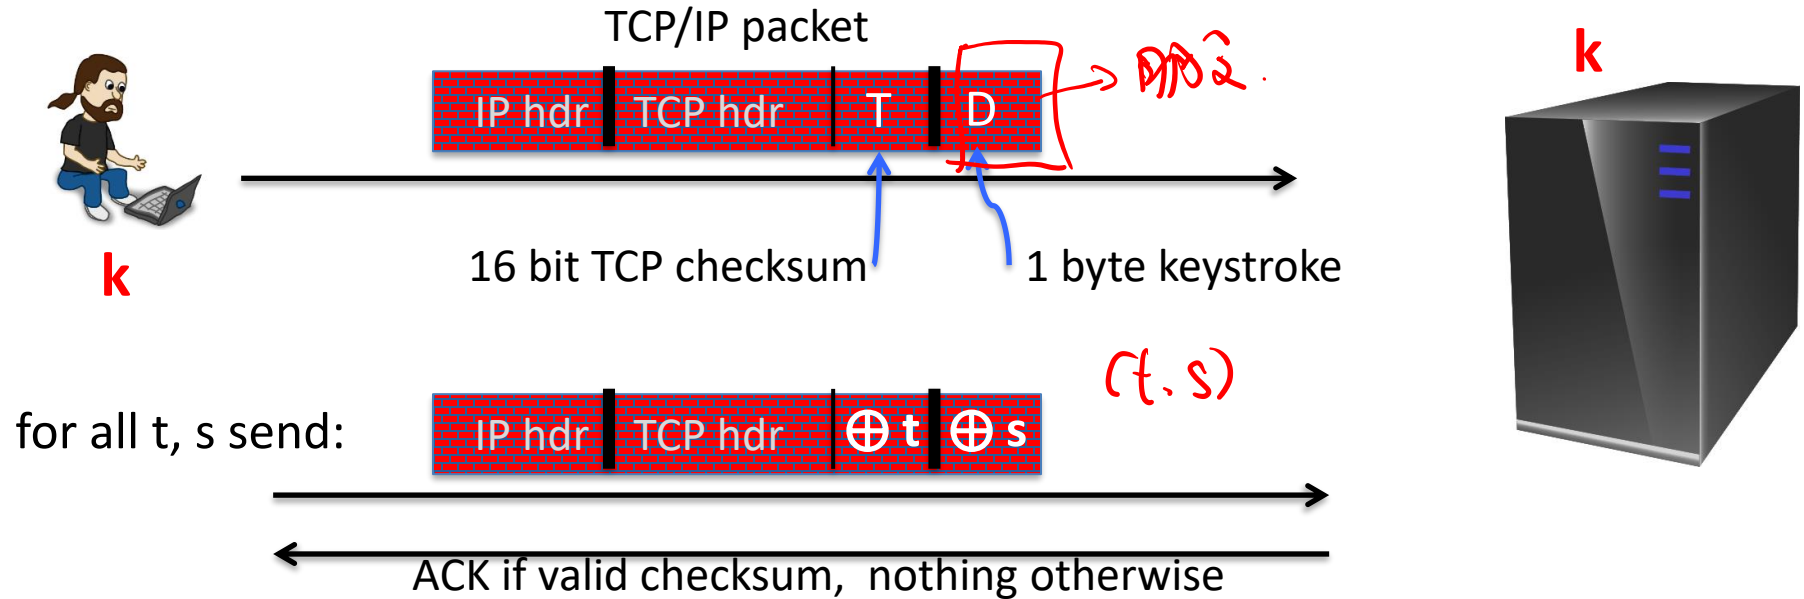
Encryption is done with CBC with a random IV.

What should IV' be?  $m[0] = D(k, c[0]) \oplus \text{IV} = \text{"dest=80..."}$

- ☐  $IV' = IV \oplus (...25...)$
- ☐  $IV' = IV \oplus (...80...)$
- ☐  $IV' = IV \oplus (...80...) \oplus (...25...)$
- ☐ It can't be done

# An attack using only network access

Remote terminal app.: each keystroke encrypted with CTR mode



$$\{ \text{checksum}(\text{hdr}, D) = t \oplus \text{checksum}(\text{hdr}, D \oplus s) \} \Rightarrow \text{can find } D$$



# The lesson

CPA security cannot guarantee secrecy under active attacks.

主动攻击

Only use one of two modes:

- If message needs integrity but no confidentiality:  
use a **MAC**
- If message needs both integrity and confidentiality:  
use **authenticated encryption** modes (this module)



# Authenticated Encryption

---

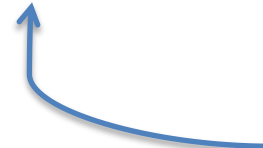
## Definitions

# Goals

An **authenticated encryption** system  $(E,D)$  is a cipher where

As usual:  $E: \underline{K} \times \underline{M} \times \cancel{N} \rightarrow \underline{C}$

but  $D: \underline{K} \times \underline{C} \times \cancel{N} \rightarrow \underline{M} \cup \{\perp\}$


 ciphertext  
is rejected

Security: the system must provide

- sem. security under a CPA attack, and
- **ciphertext integrity**:  
attacker cannot create new ciphertexts that decrypt properly  $\neq \perp$

# Ciphertext integrity

Let  $(E,D)$  be a cipher with message space  $M$ .



$$\begin{cases} b=1 & \text{if } D(k,c) \neq \perp \text{ and } c \notin \{c_1, \dots, c_q\} \\ b=0 & \text{otherwise} \end{cases}$$

Def:  $(E,D)$  has **ciphertext integrity** if for all “efficient”  $A$ :

$$\text{Adv}_{\text{CI}}[A,E] = \Pr[\text{Chal. outputs 1}] \text{ is “negligible.”}$$

# Authenticated encryption

Def: cipher  $(E,D)$  provides authenticated encryption (AE) if it is

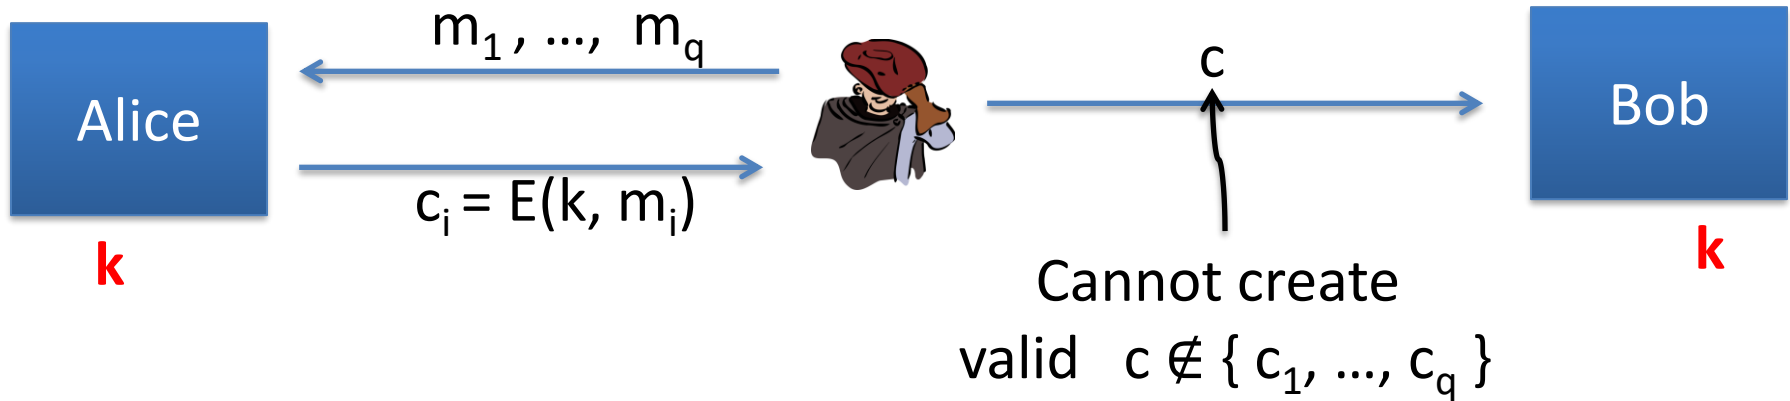
- (1) semantically secure under CPA, and
- (2) has ciphertext integrity

Bad example: CBC with rand. IV does not provide AE

- $D(k,\cdot)$  never outputs  $\perp$ , hence adv. easily wins CI game

# Implication 1: authenticity

Attacker cannot fool Bob into thinking a message was sent from Alice



$\Rightarrow$  if  $D(k, c) \neq \perp$  Bob knows message is from someone who knows  $k$   
(but message could be a replay)

# Implication 2

Authenticated encryption  $\Rightarrow$

Security against **chosen ciphertext attacks**  
(next segment)



# Authenticated Encryption

---


Chosen ciphertext  
attacks




# Example chosen ciphertext attacks

Adversary has ciphertext  $c$  that it wants to decrypt

- Often, adv. can fool server into decrypting **certain** ciphertexts (not  $c$ )



- Often, adversary can learn partial information about plaintext



# Chosen ciphertext security


**Adversary's power:** both CPA and CCA

- Can obtain the encryption of arbitrary messages of his choice
- Can decrypt any ciphertext of his choice, other than challenge  
(conservative modeling of real life)

**Adversary's goal:** Break semantic security

# Chosen ciphertext security: definition

$\mathbb{E} = (E, D)$  cipher defined over  $(K, M, C)$ . For  $b=0,1$  define  $\text{EXP}(b)$ :




# Chosen ciphertext security: definition

$\mathbb{E}$  is CCA secure if for all “efficient”  $A$ :

$$\text{Adv}_{\text{CCA}}[A, \mathbb{E}] = \left| \Pr[\text{EXP}(0)=1] - \Pr[\text{EXP}(1)=1] \right| \text{ is “negligible.”}$$

**Example:** CBC with rand. IV is not CCA-secure



# Authenticated enc. $\Rightarrow$ CCA security

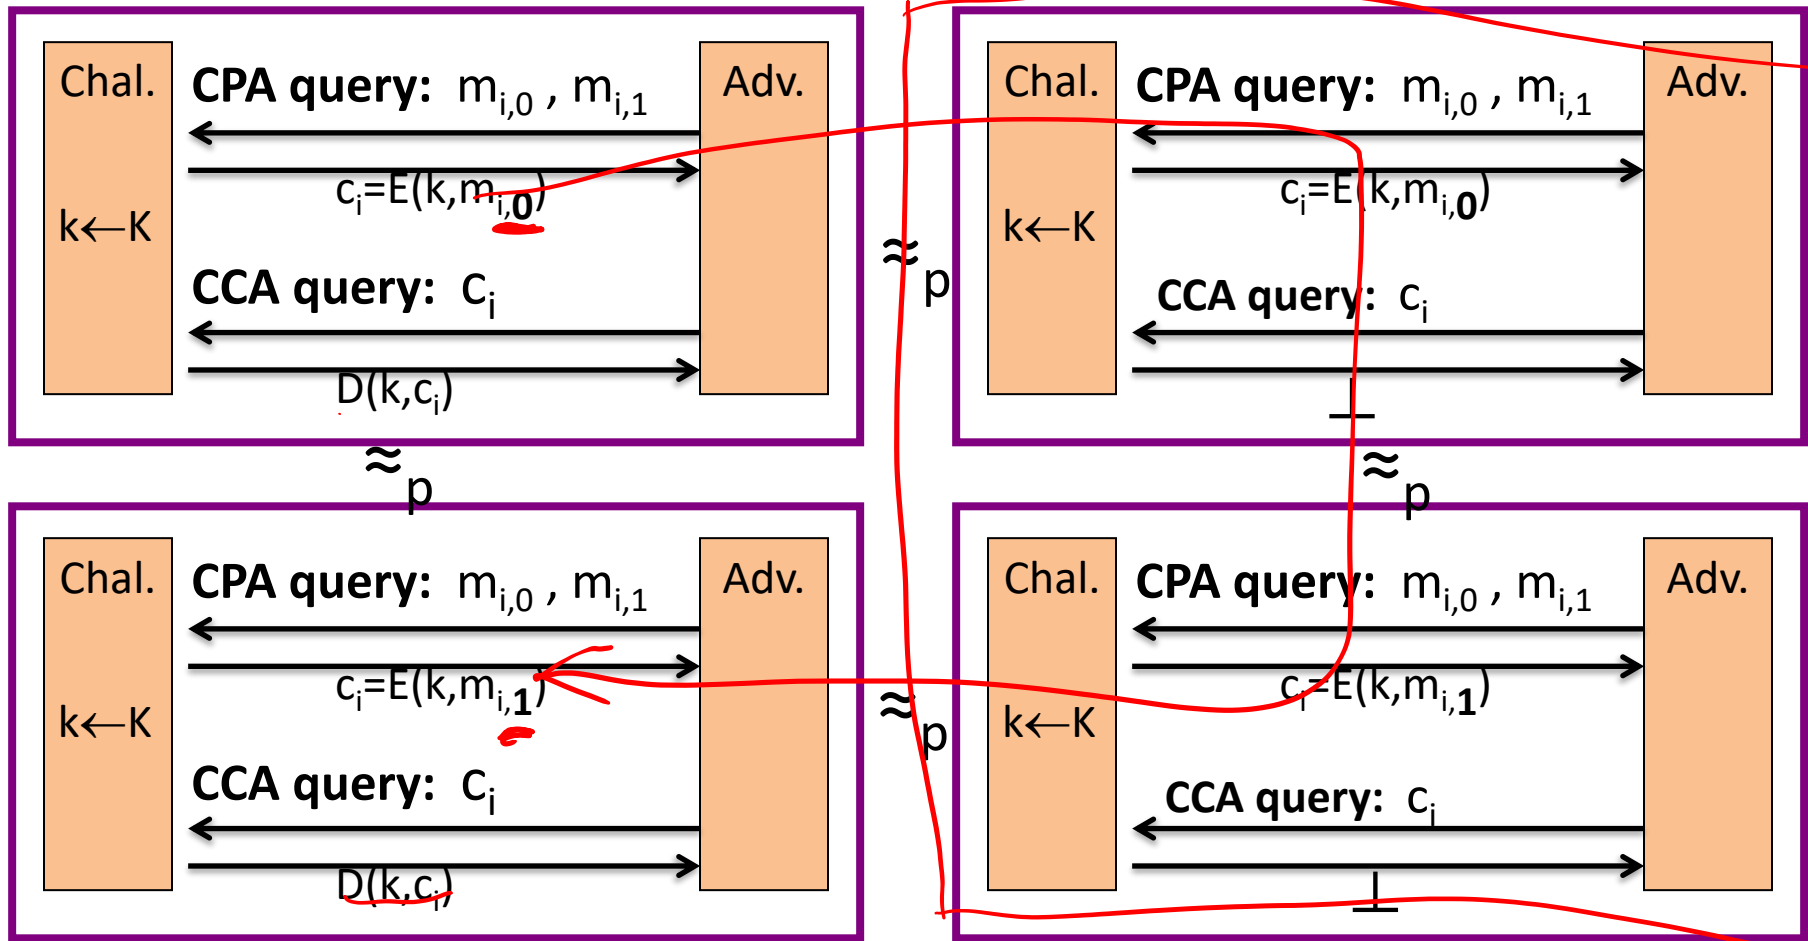
**Thm:** Let  $(E,D)$  be a cipher that provides AE.  
Then  $(E,D)$  is CCA secure !

机密 CPA, CCA  
机密完整性.

In particular, for any q-query eff. A there exist eff.  $B_1, B_2$  s.t.

$$\text{Adv}_{\text{CCA}}[A,E] \leq 2q \cdot \text{Adv}_{\text{CI}}[B_1,E] + \text{Adv}_{\text{CPA}}[B_2,E]$$

# Proof by pictures



# So what?

Authenticated encryption:

- ensures confidentiality against an active adversary that can decrypt some ciphertexts

Limitations:

- does not prevent replay attacks
- does not account for side channels (timing)



# Authenticated Encryption

---

Constructions from  
ciphers and MACs



## ... but first, some history

Authenticated Encryption (AE): introduced in 2000 [KY'00, BN'00]

Crypto APIs before then: (e.g. MS-CAPI) *crypto API*

- Provide API for CPA-secure encryption (e.g. CBC with rand. IV)
- Provide API for MAC (e.g. HMAC)

Every project had to combine the two itself without a well defined goal

- Not all combinations provide AE ...


# Combining MAC and ENC (CCA)

Encryption key  $k_E$ .

MAC key =  $k_I$

Option 1: (SSL)


*MAC - then - Enc*



Option 2: (IPsec)


*Enc - then MAC* ✓

**always  
correct**



Option 3: (SSH)

*Enc - and - MAC*



# A.E. Theorems

Let  $(E,D)$  be CPA secure cipher and  $(S,V)$  secure MAC. Then:

1. **Encrypt-then-MAC:** always provides A.E.
2. **MAC-then-encrypt:** may be insecure against CCA attacks


however: when  $(E,D)$  is rand-CTR mode or rand-CBC  
M-then-E provides A.E.

for rand-CTR mode, one-time MAC is sufficient

# Standards (at a high level)

- NIST* {
- **GCM:** CTR mode encryption then CW-MAC  
(accelerated via Intel's PCLMULQDQ instruction)
  - **CCM:** CBC-MAC then CTR mode encryption (802.11i)  
*→ AES* *→ AES*
  - **EAX:** CTR mode encryption then CMAC

All support AEAD: (auth. enc. with associated data). All are nonce-based.



# An example API (OpenSSL)

```
int AES_GCM_Init(AES_GCM_CTX *ain,  
    unsigned char *nonce, unsigned long noncelen,  
    unsigned char *key, unsigned int klen )
```


```
int AES_GCM_EncryptUpdate(AES_GCM_CTX *a,  
    unsigned char *aad, unsigned long aadlen,  
    unsigned char *data, unsigned long datalen,  
    unsigned char *out, unsigned long *outlen)
```

# MAC Security -- an explanation

Recall: MAC security implies  $(\underline{m}, t) \not\Rightarrow (\underline{m}, \textcircled{t'})$

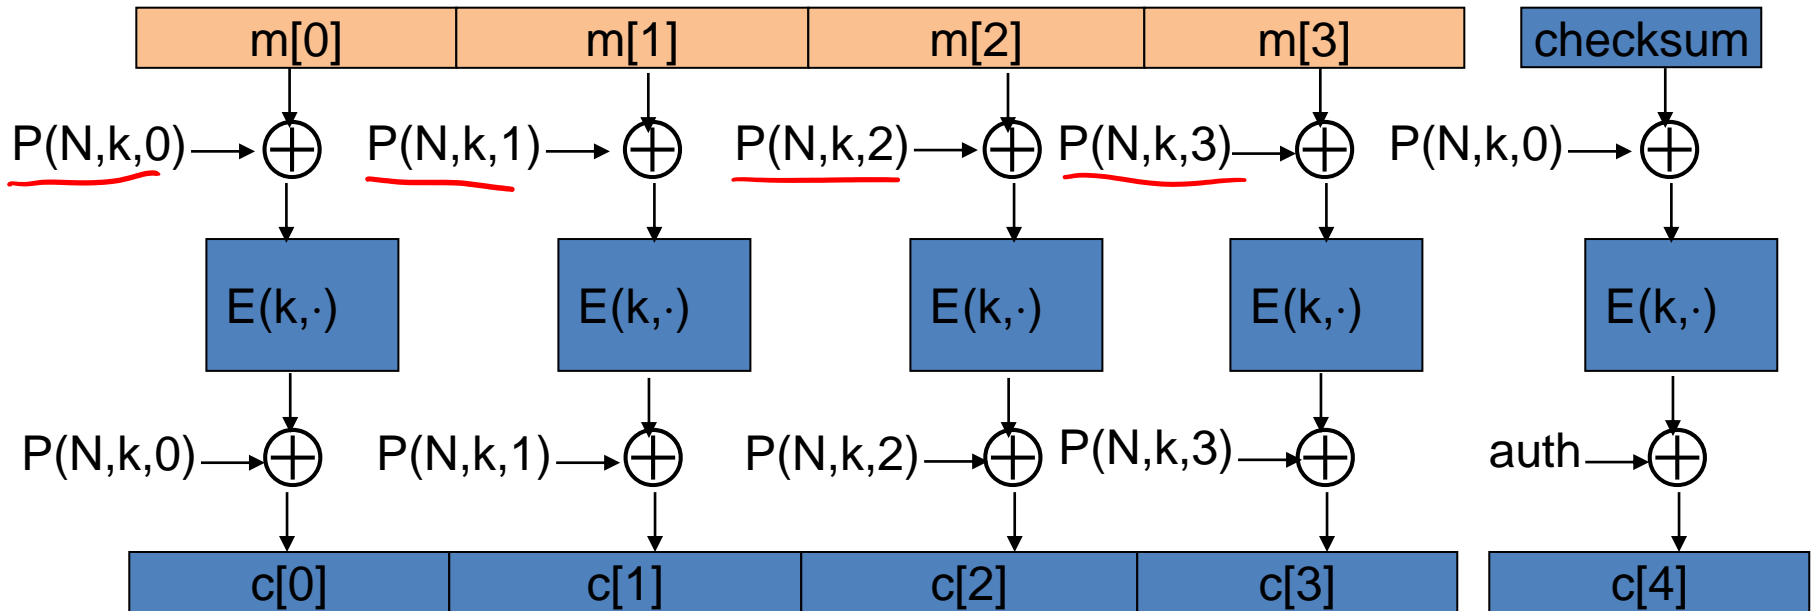
Why? Suppose not:  $(m, t) \rightarrow (m, t')$

Then Encrypt-then-MAC would not have Ciphertext Integrity !!



# OCB: a direct construction from a PRP

**More efficient authenticated encryption:** one  $E()$  op. per block.



# Performance:

Crypto++ 5.6.0 [ Wei Dai ]

AMD Opteron, 2.2 GHz (Linux)

<u>Cipher</u>	<u>code size</u>	<u>Speed (MB/sec)</u>		
AES/GCM	large**	108	AES/CTR	139
AES/CCM	smaller	61	AES/CBC	109
AES/EAX	smaller	61	AES/CMAC	109
AES/OCB		129*	HMAC/SHA1	147

\* extrapolated from Ted Kravitz's results

\*\* non-Intel machines





# Authenticated Encryption

---

Case study: TLS

# The TLS Record Protocol (TLS 1.2)



Unidirectional keys:  $k_{b \rightarrow s}$  and  $k_{s \rightarrow b}$

Stateful encryption:

- Each side maintains two 64-bit counters:  $ctr_{b \rightarrow s}$  ,  $ctr_{s \rightarrow b}$
- Init. to 0 when session started.  $ctr++$  for every record.
- Purpose: replay defense

# TLS record: encryption (CBC AES-128, HMAC-SHA1)

$$k_{b \rightarrow s} = (k_{\text{mac}}, k_{\text{enc}})$$



Browser side  $\text{enc}(k_{b \rightarrow s}, \text{data}, \text{ctr}_{b \rightarrow s})$ :

*not transmitted in packet*

step 1:  $\text{tag} \leftarrow S(k_{\text{mac}}, [ ++\text{ctr}_{b \rightarrow s} \parallel \text{header} \parallel \text{data} ])$

step 2: pad  $[ \text{header} \parallel \text{data} \parallel \text{tag} ]$  to AES block size

step 3: CBC encrypt with  $k_{\text{enc}}$  and new random IV

step 4: prepend header

# TLS record: decryption (CBC AES-128, HMAC-SHA1)

Server side **dec( $k_{b \rightarrow s}$  , record,  $ctr_{b \rightarrow s}$  ) :**

step 1: CBC decrypt record using  $k_{enc}$

step 2: check pad format: send bad\_record\_mac if invalid

step 3: check tag on [ ++ $ctr_{b \rightarrow s}$  || header || data]

send bad\_record\_mac if invalid

Provides authenticated encryption

(provided no other info. is leaked during decryption)

# Bugs in older versions (prior to TLS 1.1)

**IV for CBC is predictable:** (chained IV)

IV for next record is last ciphertext block of current record.

Not CPA secure. (a practical exploit: BEAST attack)

**Padding oracle:** during decryption

12}

if pad is invalid send **decryption failed** alert

if mac is invalid send **bad\_record\_mac** alert

⇒ attacker learns info. about plaintext (attack in next segment)

Lesson: when decryption fails, do not explain why

# Leaking the length

The TLS header leaks the length of TLS records

- Lengths can also be inferred by observing network traffic

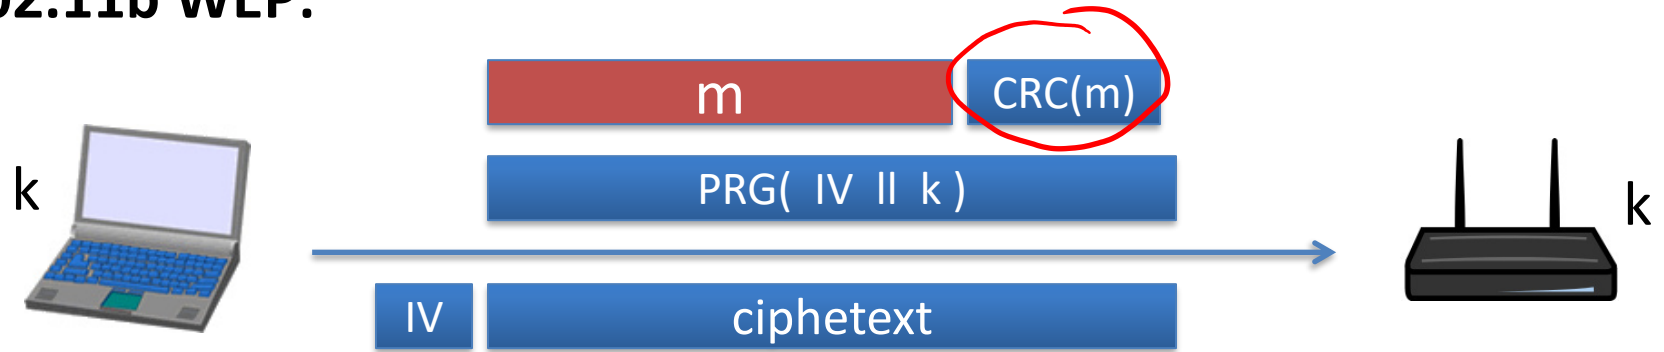
For many web applications, leaking lengths reveals sensitive info:

- In tax preparation sites, lengths indicate the type of return being filed which leaks information about the user's income
- In healthcare sites, lengths leaks what page the user is viewing
- In Google maps, lengths leaks the location being requested

No easy solution

# 802.11b WEP: how not to do it

## 802.11b WEP:




$(\text{key} || k)$

Previously discussed problems:  
two time pad and related PRG seeds

# Active attacks

**Fact:** CRC is linear, i.e.  $\forall m, p: \text{CRC}(m \oplus p) = \text{CRC}(m) \oplus \text{F}(p)$

WEP ciphertext:



attacker:

$$XX = 25 \oplus 80$$



Upon decryption: CRC is valid, but ciphertext is changed !!





# Authenticated Encryption

---

## CBC paddings attacks

# Recap

**Authenticated encryption:** CPA security + ciphertext integrity

- Confidentiality in presence of **active** adversary
- Prevents chosen-ciphertext attacks

Limitation: cannot help bad implementations ... (this segment)

Authenticated encryption modes:

- Standards: GCM, CCM, EAX
- General construction: encrypt-then-MAC

# The TLS record protocol (CBC encryption)

Decryption:  $\text{dec}(k_{b \rightarrow s}, \text{record}, \text{ctr}_{b \rightarrow s})$  :


step 1: CBC decrypt record using  $k_{\text{enc}}$

step 2: check pad format: abort if invalid

step 3: check tag on  $[++\text{ctr}_{b \rightarrow s} \parallel \text{header} \parallel \text{data}]$   
abort if invalid

Two types of error:

- padding error
- MAC error



# Padding oracle

Suppose attacker can differentiate the two errors  
(pad error, MAC error):


⇒ **Padding oracle:**

attacker submits ciphertext and learns if  
last bytes of plaintext are a valid pad

Nice example of a  
**chosen ciphertext attack**



# Padding oracle via timing OpenSSL



Credit: Brice Canvel

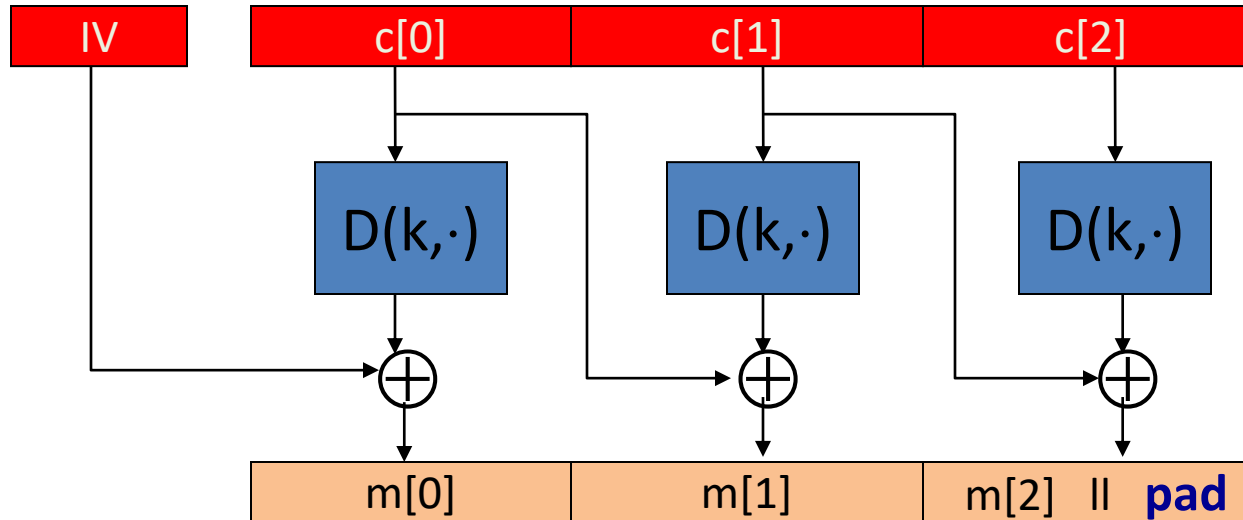
(fixed in OpenSSL 0.9.7a)

In older TLS 1.0: padding oracle due to different alert messages.

# Using a padding oracle (CBC encryption)

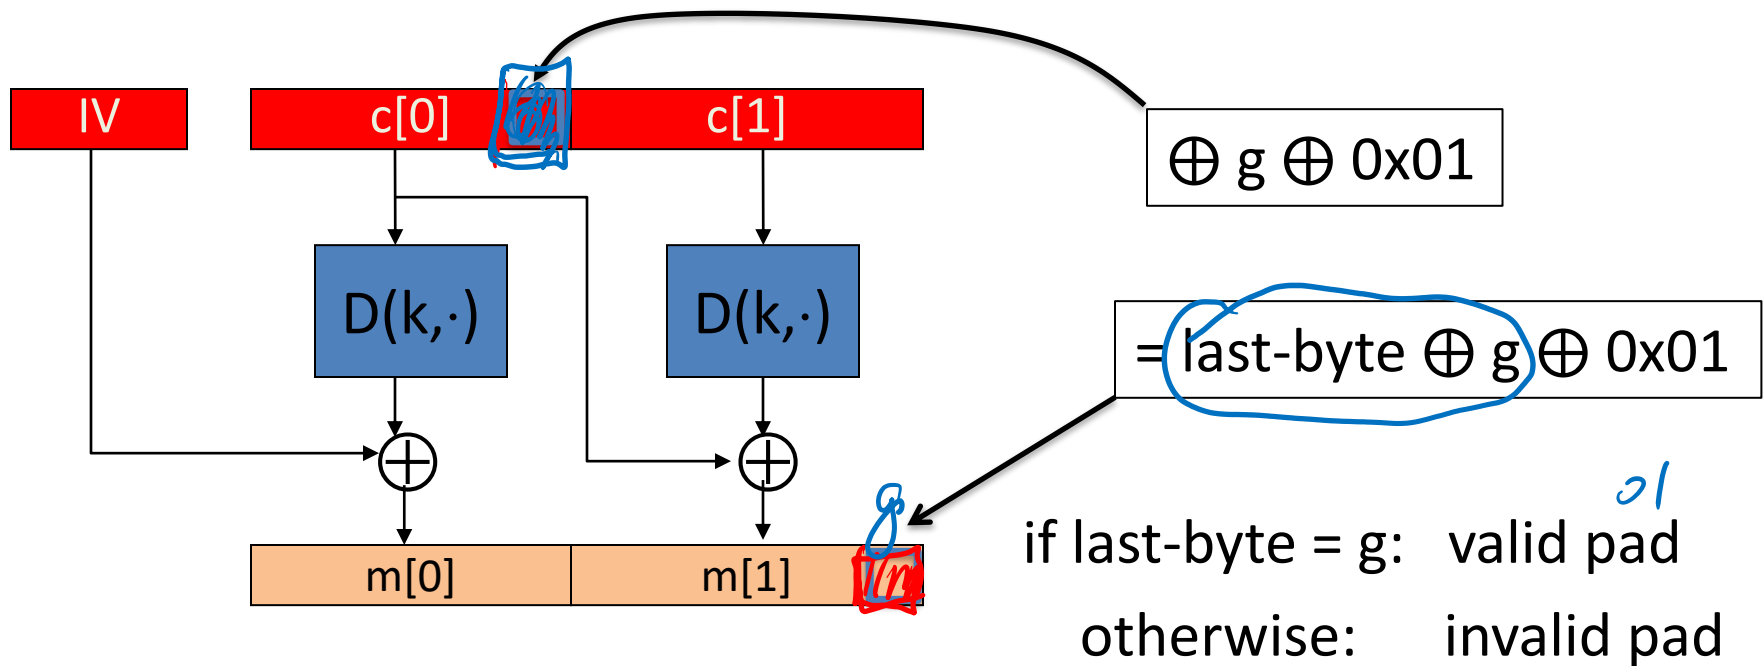
Attacker has ciphertext  $\mathbf{c} = (\mathbf{c}[0], \mathbf{c}[1], \mathbf{c}[2])$  and it wants  $\mathbf{m}[1]$

*pad error*  
*mac error*



# Using a padding oracle (CBC encryption)

step 1: let  $g$  be a guess for the last byte of  $m[1]$




# Using a padding oracle (CBC encryption)

Attack: submit  $(IV, c'[0], c[1])$  to padding oracle

$\Rightarrow$  attacker learns if last-byte =  $g$

Repeat with  $g = 0, 1, \dots, 255$  to learn last byte of  $m[1]$  01

Then use a  pad to learn the next byte and so on ...  $g_1, g_2, \oplus \underline{02, 02}$



# IMAP over TLS

**Problem:** TLS renegotiates key when an invalid record is received

Enter IMAP over TLS: (protocol for reading email)

- Every five minutes client sends login message to server:

LOGIN "username" "password"

↗

pad error

- Exact same attack works, despite new keys  
⇒ recovers password in a few hours.

↓  
mitc error

# Lesson

1. Encrypt-then-MAC would completely avoid this problem:

MAC is checked first and ciphertext discarded if invalid

2. MAC-then-CBC provides A.E., but padding oracle destroys it

Will this attack work if TLS used counter mode instead of CBC?  
(i.e. use MAC-then-CTR )

- ☐ Yes, padding oracles affect all encryption schemes
- ☐ It depends on what block cipher is used
- ☐ No, counter mode need not use padding
- ☐




# Authenticated Encryption

---

Attacking non-atomic  
decryption

# SSH Binary Packet Protocol

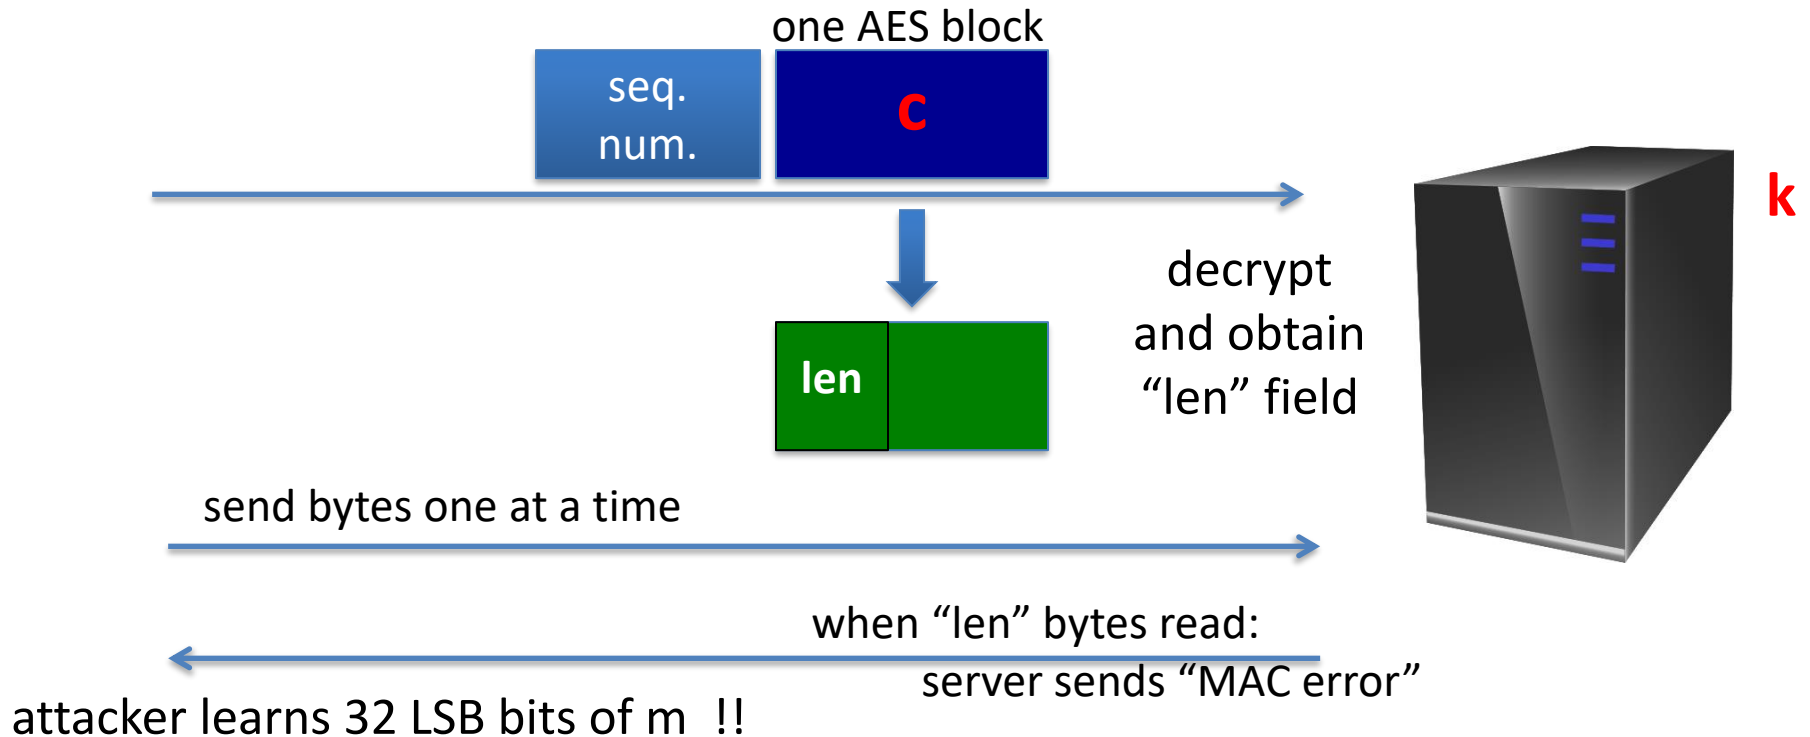


Decryption:

- step 1: decrypt packet length field only (!)
- step 2: read as many packets as length specifies
- step 3: decrypt remaining ciphertext blocks
- step 4: check MAC tag and send error response if invalid

# An attack on the enc. length field (simplified)

Attacker has one ciphertext block  $c = \text{AES}(k, m)$  and it wants  $m$



# Lesson

The problem: (1) non-atomic decrypt

(2) len field decrypted and used before it is authenticated

How would you redesign SSH to resist this attack?

- ⇒ ○ Send the length field unencrypted (but MAC-ed)
- Replace encrypt-and-MAC by encrypt-then-MAC
- ⇒ ○ Add a MAC of (seq-num, length) right after the len field
- Remove the length field and identify packet boundary by verifying the MAC after every received byte

# Further reading

- The Order of Encryption and Authentication for Protecting Communications, H. Krawczyk, Crypto 2001.
- Authenticated-Encryption with Associated-Data, P. Rogaway, Proc. of CCS 2002.
- Password Interception in a SSL/TLS Channel, B. Canvel, A. Hiltgen, S. Vaudenay, M. Vuagnoux, Crypto 2003.
- Plaintext Recovery Attacks Against SSH, M. Albrecht, K. Paterson and G. Watson, IEEE S&P 2009
- Problem areas for the IP security protocols, S. Bellare, Usenix Security 1996.